

CAIM: Cerca i Anàlisi d'Informació Massiva

FIB, Grau en Enginyeria Informàtica

Slides by Marta Arias, José Luis Balcázar,
Ramon Ferrer-i-Cancho, Ricard Gavaldá
Department of Computer Science, UPC

Fall 2018

<http://www.cs.upc.edu/~caim>

5. Web Search. Architecture of simple IR systems

Searching the Web, I

When documents are interconnected


The World Wide Web is **huge**


- ▶ 100,000 indexed pages in 1994
- ▶ 10,000,000,000's indexed pages in 2013
- ▶ Most queries will return *millions* of pages with high similarity.
- ▶ Content (text) alone cannot discriminate.
- ▶ Use the **structure** of the Web - a graph.
- ▶ Gives indications of the prestige - usefulness of each page.

How Google worked in 1998

S. Brin, L. Page: “The Anatomy of a Large-Scale Hypertextual Web Search Engine”, 1998

Notation:





Some components

- ▶ URL store: URLs awaiting exploration
- ▶ Doc repository: full documents, zipped
- ▶ Indexer: Parses pages, separates text (to Forward Index), links (to Anchors) and essential text info (to Doc Index)
 - ▶ Text in an anchor very relevant for *target* page
`anchor`
 - ▶ Font, placement in page makes some terms extra relevant
- ▶ Forward index: docid → list of terms appearing in docid
- ▶ Inverted index: term → list of docid's containing term

The inverter (sorter), I

Transforms forward index to inverted index

First idea:

```
for every entry document d  
    for every term t in d  
        add docid(d) at end of list for t;
```

Lousy locality, many disk seeks, too slow

The inverter (sorter), II

Better idea for indexing:


```
create in disk an empty inverted file, ID;  
create in RAM an empty index IR;  
for every document d  
    for every term t in d  
        add docid(d) at end of list for t in IR;  
        if RAM full  
            for each t, merge the list for t in IR  
                into the list for t in ID;
```

Merging previously sorted lists is sequential access

Much better locality. Much fewer disk seeks.

The inverter (sorter), III

The above can be done **concurrently** on different sets of documents:




The inverter (sorter), IV

- ▶ Indexer ships barrels, fragments of forward index
- ▶ Barrel size = what fits in main memory
- ▶ Separately, concurrently inverted in main memory
- ▶ Inverted barrels merged to inverted index
- ▶ 1 day instead of estimated months

Crawling


Crawler, robot, spider, wanderer ...

Systematically explores the web & collect documents.



```
add ``seed'' URLs to queue
loop
    choose a URL from queue
    fetch page, parse it
    discard it or add it to DB
    add (new) URL's it contains to queue
end loop
```

Crawling as graph exploration



Crawling process

Exploration may be:

- ▶ breadth-first, depth-first, none of the above . . .
 - ▶ focused (or not): uses expressed focus or interests
 - ▶ by keywords
 - ▶ implicitly in choice of seed pages
 - ▶ pages in the queue closer to focus get explored first
-
- ▶ Pages must be refreshed periodically.
 - ▶ Pages with higher interest fetched first, refreshed more often.

The crawling process

Crawlers must be

- ▶ efficient
- ▶ robust
- ▶ polite

Crawling efficiency

- ▶ Distributed: use several machines
- ▶ Scalable: can add more machines for more throughput
- ▶ Connections have high latency
- ▶ Keep many open connections (100's?) per machine
- ▶ Try to keep all threads busy
- ▶ DNS server tends to be the bottleneck

Crawling efficiency

Some pages may be discarded:

- ▶ Duplicates
 - ▶ Fast duplicate detection a problem in itself
 - ▶ Fingerprints or k-shingles (similar to n-grams)
- ▶ Irrelevant for crawler's goals (e.g., focused crawlers)
- ▶ Unreliable or spam

Crawling robustness

- ▶ Dead URL's: *Very common.* Timeout mechanisms
- ▶ Syntactically incorrect pages
- ▶ Spider traps. Often dynamically generated
- ▶ Webspam
- ▶ Mirror sites

Crawling politeness

- ▶ Don't hit the same server too often, esp. downloads
- ▶ Insert wait times
- ▶ Respect **robot exclusion standard**
 - ▶ /robots.txt file: administrator preferences
 - ▶ "If you are agent X, please don't explore directory Y"

```
User-agent: *
Disallow: /cgi-bin/
Disallow: /images/
Disallow: /tmp/
Disallow: /private/
```

Searching the Web, I



When documents are interconnected

The internet is **huge**

- ▶ 100,000 indexed pages in 1994
- ▶ 10,000,000,000 indexed pages at end of 2011

To find content, it is necessary to **search** for it

- ▶ We know how to deal with the **content** of the webpages
- ▶ But.. what can we do with the **structure** of the internet?



Searching the Web, II

Meaning of a hyperlink

When page A links to page B , this means

- ▶ A 's author thinks that B 's content is **interesting** or important
- ▶ So a link from A to B , adds to B 's **reputation**

But not all links are equal..

- ▶ If A is very important, then $A \rightarrow B$ “counts more”
- ▶ If A is not important, then $A \rightarrow B$ “counts less”

In today's lecture we'll see two algorithms based on this idea

- ▶ *Pagerank* (Brin and Page, oct. 98)
- ▶ *HITS* (Kleinberg, apr. 98)

Pagerank, I

The idea that made Google great

Intuition:

A page is important if it is pointed to by other important pages

- ▶ Circular definition ...
- ▶ not a problem!

Pagerank, II

Definitions

The web is a graph $G = (V, E)$

- ▶ $V = \{1, \dots, n\}$ are the nodes (that is, the pages)
- ▶ $(i, j) \in E$ if page i points to page j
- ▶ we associate to each page i , a real value p_i (i 's pagerank)
- ▶ we impose that $\sum_{i=1}^n p_i = 1$

How are the p_i 's related


- ▶ p_i depends on the values p_j of pages j pointing to i

$$p_i = \sum_{j \rightarrow i} \frac{p_j}{out(j)}$$

- ▶ where $out(j)$ is j 's outdegree

Pagerank, III

Example



$$p_i = \sum_{j \rightarrow i} \frac{p_j}{\text{out}(j)}$$

A set of $n + 1$ linear equations:

$$p_1 = \frac{p_1}{3} + \frac{p_2}{2}$$

$$p_2 = \frac{p_3}{2} + p_4$$

$$p_3 = \frac{p_1}{3}$$

$$p_4 = \frac{p_1}{3} + \frac{p_2}{2} + \frac{p_3}{2}$$

$$1 = p_1 + p_2 + p_3 + p_4$$

Whose solutions is:

$$p_1 = 6/23, p_2 = 8/23, p_3 = 2/23, p_4 = 7/23$$

Pagerank, IV

Formally

Equations

- ▶ $p_i = \sum_{j:(j,i) \in E} \frac{p_j}{\text{out}(j)}$ for each $i \in V$
- ▶ $\sum_{i=1}^n p_i = 1$

where $\text{out}(i) = |\{j : (i, j) \in E\}|$ is the *outdegree* of node i


If $|V| = n$

- ▶ $n + 1$ equations
- ▶ n unknowns

Could be solved, for example, using Gaussian elimination in time $O(n^3)$

Pagerank, V

Example, revisited



A set of linear equations:

$$\begin{pmatrix} p_1 \\ p_2 \\ p_3 \\ p_4 \end{pmatrix} = \begin{pmatrix} \frac{1}{3} & \frac{1}{2} & 0 & 0 \\ 0 & 0 & \frac{1}{2} & 1 \\ \frac{1}{3} & 0 & 0 & 0 \\ \frac{1}{3} & \frac{1}{2} & \frac{1}{2} & 0 \end{pmatrix} \cdot \begin{pmatrix} p_1 \\ p_2 \\ p_3 \\ p_4 \end{pmatrix}$$

namely: $\vec{p} = M^T \vec{p}$ and additionally
 $\sum_i p_i = 1$


Whose solutions is:

\vec{p} is the eigenvector of matrix M^T associated to eigenvalue 1

Pagerank, VI

Example, revisited

What does M^T look like?




$$M^T = \begin{pmatrix} \frac{1}{3} & \frac{1}{2} & 0 & 0 \\ 0 & 0 & \frac{1}{2} & 1 \\ \frac{1}{3} & 0 & 0 & 0 \\ \frac{1}{3} & \frac{1}{2} & \frac{1}{2} & 0 \end{pmatrix}$$

M^T is the *transpose* of the row-normalized **adjacency matrix** of the graph !

Pagerank, VII

Example, revisited

Adjacency matrix



$$A = \begin{pmatrix} 1 & 0 & 1 & 1 \\ 1 & 0 & 0 & 1 \\ 0 & 1 & 0 & 1 \\ 0 & 1 & 0 & 0 \end{pmatrix}$$

$$M = \begin{pmatrix} 1/3 & 0 & 1/3 & 1/3 \\ 1/2 & 0 & 0 & 1/2 \\ 0 & 1/2 & 0 & 1/2 \\ 0 & 1 & 0 & 0 \end{pmatrix}$$

(rows add up to 1)

$$M^T = \begin{pmatrix} 1/3 & 1/2 & 0 & 0 \\ 0 & 0 & 1/2 & 1 \\ 1/3 & 0 & 0 & 0 \\ 1/3 & 1/2 & 1/2 & 0 \end{pmatrix}$$

(columns add up to 1)

Pagerank, VIII

Example, revisited

$$A = \begin{pmatrix} 1 & 0 & 1 & 1 \\ 1 & 0 & 0 & 1 \\ 0 & 1 & 0 & 1 \\ 0 & 1 & 0 & 0 \end{pmatrix} \quad M = \begin{pmatrix} \frac{1}{3} & 0 & \frac{1}{3} & \frac{1}{3} \\ \frac{1}{2} & 0 & 0 & \frac{1}{2} \\ 0 & \frac{1}{2} & 0 & \frac{1}{2} \\ 0 & 1 & 0 & 0 \end{pmatrix} \quad M^T = \begin{pmatrix} \frac{1}{3} & \frac{1}{2} & 0 & 0 \\ 0 & 0 & \frac{1}{2} & 1 \\ \frac{1}{3} & 0 & 0 & 0 \\ \frac{1}{3} & \frac{1}{2} & \frac{1}{2} & 0 \end{pmatrix}$$

Question:

Why do we need to *row-normalize* and *transpose A*?

Answer:

- ▶ *Row normalization:* because $p_i = \sum_{j:(j,i) \in E} \frac{p_j}{\text{out}(j)}$
- ▶ *Transpose:* because $p_i = \sum_{j:(j,i) \in E} \frac{p_j}{\text{out}(j)}$, that is,
 p_i depends on i 's *incoming edges*



Pagerank, IX

It is just about solving a system of linear equations!

.. but

- ▶ How do we know a solution exists?
- ▶ How do we know it has a **single** solution?
- ▶ How can we compute it efficiently?

For example, the graph on the left has no solution.. (check it!)
but the one on the right does



Pagerank, X

How do we know a solution exists?

Luckily, we have some results from *linear algebra*

Definition

A matrix M is stochastic, if

- ▶ All entries are in the range $[0, 1]$
- ▶ Each row adds up to 1 (i.e., M is row normalized)

Theorem (Perron-Frobenius)


If M is stochastic, then it has at least one stationary vector, i.e., one non-zero vector p such that

$$M^T p = p.$$

Pagerank, XI

Equivalently: the random surfer view

Now assume M is the **transition probability matrix** between states in G



$$M = \begin{pmatrix} 1/3 & 0 & 1/3 & 1/3 \\ 1/2 & 0 & 0 & 1/2 \\ 0 & 1/2 & 0 & 1/2 \\ 0 & 1 & 0 & 0 \end{pmatrix}$$

Let $\vec{p}(t)$ be the probability over states at time t

- ▶ E.g., $p_j(0)$ is the probability of being at state j at time 0

Random surfer jumps from page i to page j with probability m_{ij}

- ▶ E.g., probability of transitioning from state 2 to state 4 is $m_{24} = 1/2$

Pagerank, XII

The random surfer view


- ▶ Surfer starts at random page according to probability distribution $\vec{p}(0)$
- ▶ At time $t > 0$, random surfer follows one of current page's links uniformly at random

$$\vec{p}(t) := M^T \vec{p}(t - 1)$$

- ▶ In the limit $t \rightarrow \infty$:
 - ▶ $\vec{p}(t) = \vec{p}(t + 1) = \vec{p}(t + 2) = \dots = \vec{p}$
 - ▶ so $\vec{p}(t) = M^T \vec{p}(t - 1)$
 - ▶ $\vec{p}(t)$ converges to a solution p s.t. $p = M^T p$ (the pagerank solution)!

Pagerank, XIII

Random surfer example



$$M^T = \begin{pmatrix} \frac{1}{3} & \frac{1}{2} & 0 & 0 \\ 0 & 0 & \frac{1}{2} & 1 \\ \frac{1}{3} & 0 & 0 & 0 \\ \frac{1}{3} & \frac{1}{2} & \frac{1}{2} & 0 \end{pmatrix}$$

- ▶ $\vec{p}(0)^T = (1, 0, 0, 0)$
- ▶ $\vec{p}(1)^T = (1/3, 0, 1/3, 1/3)$
- ▶ $\vec{p}(2)^T = (0.11, 0.50, 0.11, 0.28)$
- ▶ ..
- ▶ $\vec{p}(10)^T = (0.26, 0.35, 0.09, 0.30)$
- ▶ $\vec{p}(11)^T = (0.26, 0.35, 0.09, 0.30)$

Pagerank, XIV

An algorithm to solve the eigenvector problem (find p s.t. $p = M^T p$)

The Power Method


- ▶ Chose initial vector $\vec{p}(0)$ randomly
- ▶ Repeat $\vec{p}(t) \leftarrow M^T \vec{p}(t - 1)$
- ▶ Until convergence (i.e. $\vec{p}(t) \approx \vec{p}(t - 1)$)

We are hoping that

- ▶ The method converges
- ▶ The method converges **fast**
- ▶ The method converges fast to the **pagerank solution**
- ▶ The method converges fast to the pagerank solution
regardless of the initial vector

Pagerank, XV

Convergence of the Power method: aperiodicity required



Try out the power method with $\vec{p}(0)$:

$$\begin{pmatrix} 1/4 \\ 1/4 \\ 1/4 \\ 1/4 \end{pmatrix}, \text{ or } \begin{pmatrix} 1 \\ 0 \\ 0 \\ 0 \end{pmatrix}, \text{ or } \begin{pmatrix} 1/2 \\ 0 \\ 1/2 \\ 0 \end{pmatrix}$$


Not being able to break the **cycle** looks problematic!

- ▶ .. so will require graphs to be **aperiodic**
 - ▶ no integer $k > 1$ dividing the length of every cycle

Pagerank, XVI

Convergence of the Power method: strong connectedness required

What happens with the pagerank in this graph?



The **sink** hoards all the pagerank!

- ▶ need a way to leave sinks
- ▶ .. so we will force graphs to be **strongly connected**

Pagerank, XVII

A useful theorem from Markov chain theory

Theorem

If a matrix M is **strongly connected** and **aperiodic**, then:

- ▶ $M^T \vec{p} = \vec{p}$ has exactly one non-zero solution such that $\sum_i p_i = 1$
- ▶ 1 is the largest eigenvalue of M^T
- ▶ the Power method converges to the \vec{p} satisfying $M^T \vec{p} = \vec{p}$, from any initial non-zero $\vec{p}(0)$
- ▶ Furthermore, we have exponential fast convergence

To guarantee a solution, we will make sure that the matrices that we work with are **strongly connected** and **aperiodic**

Pagerank, XVIII

Guaranteeing aperiodicity and strong connectedness

Definition (The Google Matrix)

Given a **damping factor** λ such that: $0 < \lambda < 1$:

$$G = \lambda M + (1 - \lambda) \frac{1}{n} J$$

where J is a $n \times n$ matrix containing 1 in each entry

Observe that:

- ▶ G is stochastic
 - ▶ .. because G is a weighted average of M and $\frac{1}{n}J$, which are also stochastic
- ▶ for each integer $k > 0$, there is a non-zero probability path of length k from every state to any other state of G
 - ▶ .. implying that G is strongly connected and aperiodic
- ▶ and so the Power method will converge on G , and fast!

Pagerank, XIX

Teleportation in the random surfer view


The meaning of λ

- ▶ With probability λ , the random surfer follows a link in current page
- ▶ With probability $1 - \lambda$, the random surfer jumps to a random page in the graph (**teleportation**)

Pagerank, XX

Excercise, I

Compute the pagerank value of each node of the following graph assuming a damping factor $\lambda = 2/3$:



Hint: solve the following system, using $p_2 = p_3 = p_4$

$$\begin{pmatrix} p_1 \\ p_2 \\ p_3 \\ p_4 \end{pmatrix} = \left[\frac{2}{3} \begin{pmatrix} 0 & 1 & 1 & 1 \\ \frac{1}{3} & 0 & 0 & 0 \\ \frac{1}{3} & 0 & 0 & 0 \\ \frac{1}{3} & 0 & 0 & 0 \end{pmatrix} + \frac{1}{3} \cdot \frac{1}{4} \begin{pmatrix} 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 \end{pmatrix} \right] \begin{pmatrix} p_1 \\ p_2 \\ p_3 \\ p_4 \end{pmatrix}$$

Pagerank, XXI

Exercise, II

Compute the pagerank vector \vec{p} of graph with row-normalized matrix M for damping factor λ in closed matrix form.

Answer:

$$\vec{p} = (I - \lambda M^T)^{-1} \begin{pmatrix} \frac{1-\lambda}{n} \\ \vdots \\ \frac{1-\lambda}{n} \end{pmatrix}$$

Topic-sensitive Pagerank, I

Observe that pageranks are **independent** of user's query

- ▶ Advantages
 - ▶ Computed off-line
 - ▶ Collective reputation
- ▶ Disadvantages
 - ▶ Insensitive to particular user's needs

Topic-sensitive Pagerank, II

Assume there is a small set of K topics (sports, science, politics, ...)

- ▶ Each topic $k \in \{1, \dots, K\}$ is defined by a subset of the web pages T_k
- ▶ For each k , compute pagerank of node i for topic k :

$p_{i,k}$ = “pagerank of node i with teleportation reduced to T_k ”

- ▶ Finally compute ranking score of a page i given query q

$$score(i, q) = \sum_{k=1}^K sim(T_k, q) \cdot p_{i,k}$$

HITS, I

Hypertext Induced Text Search

Interest of a web page due to two different reasons

- ▶ page **content** is interesting (*authority*), or
- ▶ page **points to** interesting pages (*hub*)

HITS main rationale

- ▶ hubs are important if they point to important authorities
- ▶ authorities are important if pointed to by important hubs
- ▶ .. but .. circular definition again **not a problem!**

HITS, II

Definition of authority and hub value (a_i and h_i)

Associate to each page i an authority value a_i and a hub value h_i

- ▶ vector of all authority values is \vec{a}
- ▶ vector of all hub values is \vec{h}

Keep these vectors normalized (notice L2 norm!)

- ▶ $\|\vec{a}\| = \sum_i a_i^2 = 1$, and $\|\vec{h}\| = \sum_i h_i^2 = 1$

For appropriate scaling constants c and d

- ▶ $a_i = c \cdot \sum_{j \rightarrow i} h_j$, and $h_i = d \cdot \sum_{i \rightarrow j} a_j$


Notice not a linear system anymore!

- ▶ ... but still ok with a variant of the power method

HITS, III

Example

Our old graph



Adjacency matrix

$$A = \begin{pmatrix} 1 & 0 & 1 & 1 \\ 1 & 0 & 0 & 1 \\ 0 & 1 & 0 & 1 \\ 0 & 1 & 0 & 0 \end{pmatrix}$$


$$a_1 = c \cdot (h_1 + h_2) \quad // \text{here we use } A \text{'s first column}$$

$$a_1 \propto (1, 1, 0, 0) \cdot \begin{pmatrix} h_1 \\ h_2 \\ h_3 \\ h_4 \end{pmatrix} = (1, 1, 0, 0) \cdot \vec{h}$$

HITS, IV

Example

Our old graph



Adjacency matrix

$$A = \begin{pmatrix} 1 & 0 & 1 & 1 \\ 1 & 0 & 0 & 1 \\ 0 & 1 & 0 & 1 \\ 0 & 1 & 0 & 0 \end{pmatrix}$$

$$h_2 = d \cdot (a_1 + a_4) \quad // \text{here we use } A \text{'s second row}$$

$$h_2 \propto (1, 0, 0, 1) \cdot \begin{pmatrix} a_1 \\ a_2 \\ a_3 \\ a_4 \end{pmatrix} = (1, 0, 0, 1) \cdot \vec{a}$$

HITS, V

Update rule for \vec{a} and \vec{h}

Written in compact matrix form

- ▶ To update authority values
 - ▶ $\vec{a} := A^T \cdot \vec{h}$
 - ▶ normalize afterwards $\vec{a} := \frac{\vec{a}}{\|\vec{a}\|}$ so that $\|\vec{a}\| = 1$
- ▶ To update hub values
 - ▶ $\vec{h} := A \cdot \vec{a}$
 - ▶ normalize afterwards $\vec{h} := \frac{\vec{h}}{\|\vec{h}\|}$ so that $\|\vec{h}\| = 1$

HITS, VI

The power method for finding \vec{a} and \vec{h}

Given adjacency matrix A

- ▶ Initialize $\vec{a} = \vec{h} = (1, 1, \dots, 1)^T$
- ▶ Normalize \vec{a} and \vec{h} so that $\|a\| = \|h\| = 1$
- ▶ Repeat until convergence
 - ▶ $\vec{a} := A^T \cdot \vec{h}$
 - ▶ normalize \vec{a} so that $\|a\| = 1$
 - ▶ $\vec{h} := A \cdot \vec{a}$
 - ▶ normalize \vec{h} so that $\|h\| = 1$

HITS, VII

HITS algorithm

Query answering algorithm HITS

- ▶ Get query q and run content-based searcher on q
- ▶ Let RootSet be the top- k ranked pages
- ▶ Expand pages to BaseSet by adding all pages pointed to and by pages in RootSet
- ▶ compute hub and authority values for the subgraph of web induced by BaseSet
- ▶ Rank pages in BaseSet according to \vec{a} , \vec{h} , and content

HITS, VIII

HITS algorithm illustrated




FIG. 1. Expanding the root set into a base set.

HITS vs. Pagerank

Pros of HITS vs. Pagerank

- ▶ Sensitive to user queries

Cons of HITS vs. Pagerank

- ▶ Compute online, not offline!
- ▶ More vulnerable to webspamming