


PARALLELISM LABORATORY 5

David Morais

david.morais

Laura van Dinteren

laura.van.dinteren



| | | | | | | |
|---|--|---------|--------|--------|--------|----------|
| | Number of threads (elapsed in seconds) | | | | | |
| Version | 1 | 4 | 8 | 12 | 16 | 20 |
| 1D Block Geometric Data Decomposition by columns | 2.86 | 1.80 | 1.25 | 0.91 | 0.73 | 0.61 |
| 1D Block-Cyclic Geometric Data Decomposition by columns | 2.858 | 0.7212 | 0.338 | 0.2691 | 0.2095 | 0.16 |
| 1D Cyclic Geometric Data Decomposition by rows | 2.85 | 0.7620 | 0.3837 | 0.2644 | 0.1994 | 0.160208 |
| | Number of threads (L2 Cache Misses per thread) | | | | | |
| 1D Block Geometric Data Decomposition by columns | 184454 | 93139 | 93225 | 93064 | 93407 | 93284 |
| 1D Block-Cyclic Geometric Data Decomposition by columns | 1642609 | 1358022 | 206380 | 138037 | 103568 | 83112 |
| 1D Cyclic Geometric Data Decomposition by rows | 1642385 | 412451 | 206881 | 138530 | 103856 | 83762 |

1D Block Geometric Data Decomposition by columns

Modelfactor tables

| Overview of whole program execution metrics | | | | | | | | | | | |
|---|------|------|------|------|------|------|------|------|------|------|------|
| Number of processors | 1 | 2 | 4 | 6 | 8 | 10 | 12 | 14 | 16 | 18 | 20 |
| Elapsed time (sec) | 2.88 | 2.09 | 1.83 | 1.58 | 1.28 | 1.10 | 0.94 | 0.86 | 0.76 | 0.71 | 0.64 |
| Speedup | 1.00 | 1.38 | 1.58 | 1.82 | 2.25 | 2.62 | 3.08 | 3.34 | 3.77 | 4.04 | 4.49 |
| Efficiency | 1.00 | 0.69 | 0.39 | 0.30 | 0.28 | 0.26 | 0.26 | 0.24 | 0.24 | 0.22 | 0.22 |

Table 1: Analysis done on Wed Dec 11 04:45:46 PM CET 2024, par4115

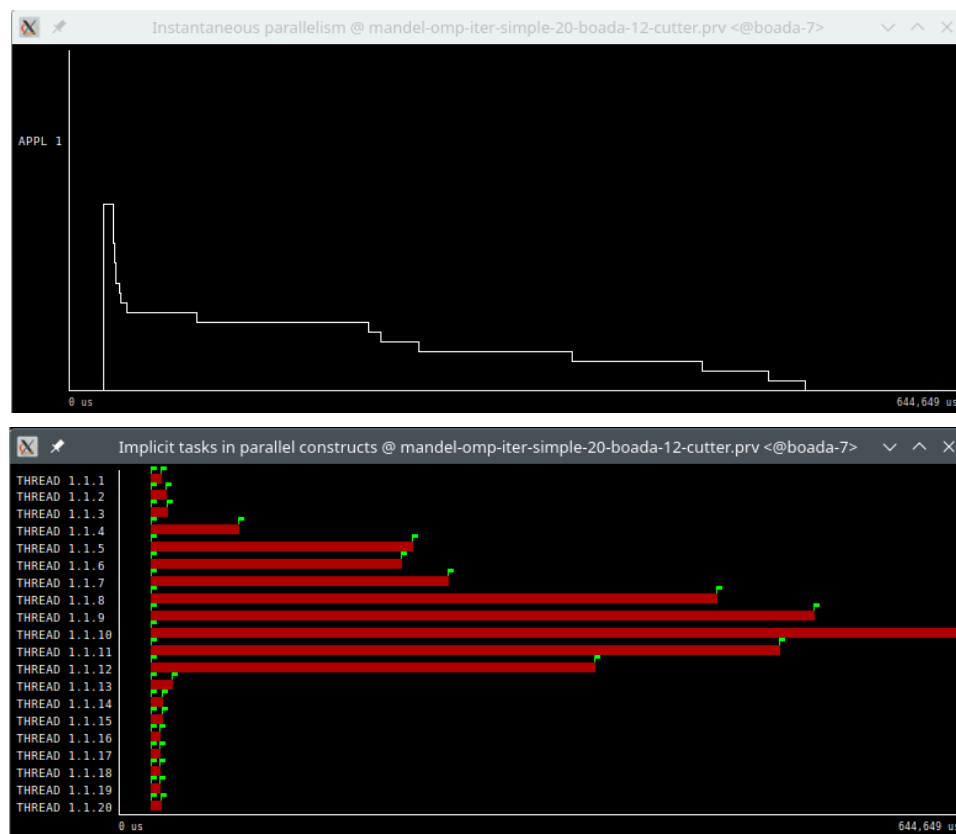
| Overview of the Efficiency metrics in parallel fraction, $\phi=99.10\%$ | | | | | | | | | | | |
|---|---------|---------|---------|---------|---------|---------|---------|--------|--------|--------|--------|
| Number of processors | 1 | 2 | 4 | 6 | 8 | 10 | 12 | 14 | 16 | 18 | 20 |
| Global efficiency | 100.00% | 69.08% | 39.60% | 30.60% | 28.50% | 26.55% | 26.15% | 24.42% | 24.16% | 23.05% | 23.19% |
| Parallelization strategy efficiency | 100.00% | 69.16% | 39.71% | 31.31% | 29.76% | 28.22% | 27.96% | 26.46% | 26.43% | 25.43% | 25.72% |
| Load balancing | 100.00% | 69.16% | 39.72% | 31.32% | 29.78% | 28.25% | 28.00% | 26.50% | 26.48% | 25.49% | 25.80% |
| In execution efficiency | 100.00% | 99.99% | 99.96% | 99.95% | 99.93% | 99.90% | 99.85% | 99.84% | 99.81% | 99.77% | 99.70% |
| Scalability for computation tasks | 100.00% | 99.88% | 99.73% | 97.72% | 95.76% | 94.06% | 93.53% | 92.28% | 91.39% | 90.64% | 90.15% |
| IPC scalability | 100.00% | 100.01% | 100.00% | 100.04% | 100.04% | 100.05% | 100.06% | 99.91% | 99.90% | 99.77% | 99.73% |
| Instruction scalability | 100.00% | 100.00% | 100.00% | 100.00% | 100.00% | 99.99% | 99.99% | 99.99% | 99.99% | 99.99% | 99.99% |
| Frequency scalability | 100.00% | 99.87% | 99.73% | 97.69% | 95.73% | 94.02% | 93.47% | 92.37% | 91.49% | 90.85% | 90.41% |

Table 2: Analysis done on Wed Dec 11 04:45:46 PM CET 2024, par4115

| Statistics about explicit tasks in parallel fraction | | | | | | | | | | | |
|--|------------|------------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|
| Number of processors | 1 | 2 | 4 | 6 | 8 | 10 | 12 | 14 | 16 | 18 | 20 |
| Number of implicit tasks per thread (average us) | 1.0 | 1.0 | 1.0 | 1.0 | 1.0 | 1.0 | 1.0 | 1.0 | 1.0 | 1.0 | 1.0 |
| Useful duration for implicit tasks (average us) | 2858756.12 | 1431055.85 | 716647.01 | 487559.46 | 373151.23 | 303918.19 | 254713.15 | 221275.21 | 195496.16 | 175221.67 | 158550.03 |
| Load balancing for implicit tasks | 1.0 | 0.69 | 0.4 | 0.31 | 0.3 | 0.28 | 0.28 | 0.27 | 0.26 | 0.25 | 0.26 |
| Time in synchronization implicit tasks (average us) | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| Time in fork/join implicit tasks (average us) | 31.7 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |

Table 3: Analysis done on Wed Dec 11 04:45:46 PM CET 2024, par4115

Paraver analysis




Memory analysis

The figure shows a screenshot of a memory analysis window titled "l2_data_cache_misses in i...12-cutter.prv <@boada-7>". The window displays a table of L2 data cache misses for various threads. The table has two columns: the thread identifier and the number of misses. The threads are listed from THREAD 1.1.1 to THREAD 1.1.20. The total number of misses is 9,999,999,999,983,222,784.00. The table also includes summary statistics: Total (2,970,217), Average (148,510.85), Maximum (185,889), Minimum (93,064), StDev (33,188.52), and Avg/Max (0.80).

| Thread | Misses |
|----------------|-------------------|
| THREAD 1.1.1 | 184,454 |
| THREAD 1.1.2 | 161,294 |
| THREAD 1.1.3 | 161,153 |
| THREAD 1.1.4 | 93,139 |
| THREAD 1.1.5 | 180,822 |
| THREAD 1.1.6 | 159,414 |
| THREAD 1.1.7 | 157,888 |
| THREAD 1.1.8 | 93,225 |
| THREAD 1.1.9 | 169,692 |
| THREAD 1.1.10 | 154,359 |
| THREAD 1.1.11 | 157,070 |
| THREAD 1.1.12 | 93,064 |
| THREAD 1.1.13 | 182,517 |
| THREAD 1.1.14 | 162,361 |
| THREAD 1.1.15 | 161,975 |
| THREAD 1.1.16 | 93,407 |
| THREAD 1.1.17 | 185,889 |
| THREAD 1.1.18 | 162,913 |
| THREAD 1.1.19 | 162,297 |
| THREAD 1.1.20 | 93,284 |
| Total | 2,970,217 |
| Average | 148,510.85 |
| Maximum | 185,889 |
| Minimum | 93,064 |
| StDev | 33,188.52 |
| Avg/Max | 0.80 |

Strong scalability



par4115
Speed-up wrt sequential time (mandel function only)
Generated by par4115 on Wed Dec 11 04:49:07 PM CET 2024

Analysis

If we take a look at the modelfactor tables we can quickly see that the block geometric data decomposition is not a good strategy for 20 threads, the speed-up is not even 5, the efficiency is 22%, and the load balancing is really poor (26%). The paraver graphics back this theory, the load balancing is awful, as a result the instantaneous parallelism decreases rapidly as the execution goes. It is also noticeable that the graphic matches the Mandelbrot computing areas. The areas that are white require more computation and the threads that execute said areas have more work to do, this explains the load unbalance. About the memory analysis, threads with lower misses might be benefiting from better spatial or temporal locality of data, whereas threads with high misses may be accessing data in less cache-friendly patterns. Finally, for the strong scalability graphic, we can see that it is lineal and it is approximately 5 for 20 threads, which is really poor.

1D Block-Cyclic Geometric Data Decomposition by columns

Modelfactor tables

| Overview of whole program execution metrics | | | | | | | | | | | |
|---|------|------|------|------|------|------|------|-------|-------|-------|-------|
| Number of processors | 1 | 2 | 4 | 6 | 8 | 10 | 12 | 14 | 16 | 18 | 20 |
| Elapsed time (sec) | 2.88 | 1.46 | 0.75 | 0.54 | 0.42 | 0.35 | 0.30 | 0.26 | 0.24 | 0.21 | 0.20 |
| Speedup | 1.00 | 1.97 | 3.86 | 5.33 | 6.82 | 8.33 | 9.73 | 11.25 | 12.13 | 13.64 | 14.70 |
| Efficiency | 1.00 | 0.99 | 0.97 | 0.89 | 0.85 | 0.83 | 0.81 | 0.80 | 0.76 | 0.76 | 0.74 |

Table 1: Analysis done on Wed Dec 18 04:24:29 PM CET 2024, par4115

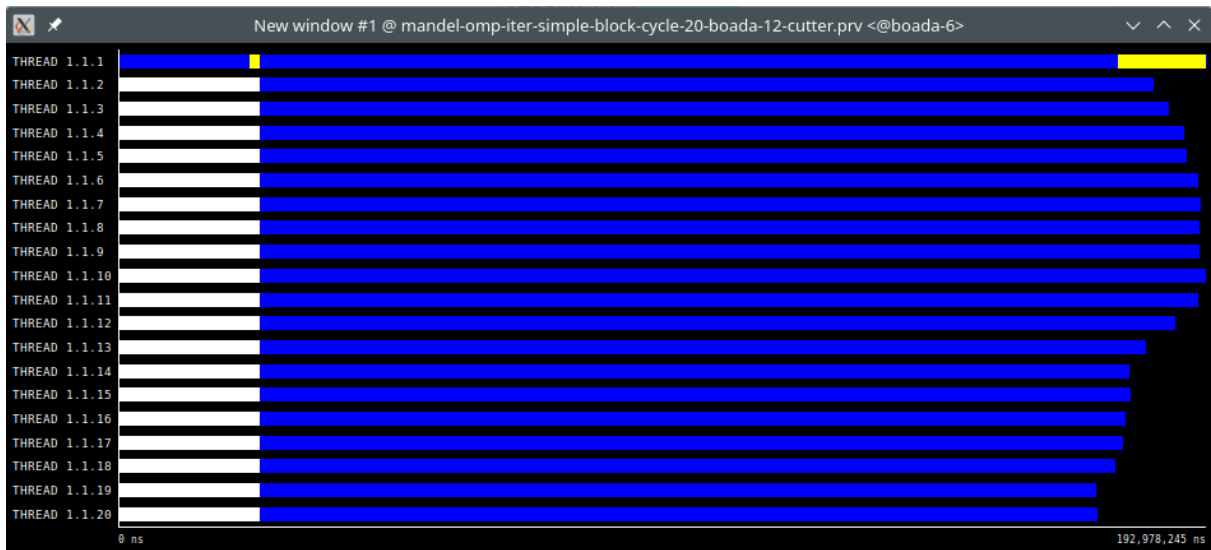
| Overview of the Efficiency metrics in parallel fraction, $\phi=99.14\%$ | | | | | | | | | | | |
|---|---------|---------|---------|---------|---------|--------|--------|--------|--------|--------|--------|
| Number of processors | 1 | 2 | 4 | 6 | 8 | 10 | 12 | 14 | 16 | 18 | 20 |
| Global efficiency | 100.00% | 99.60% | 98.88% | 92.46% | 89.74% | 89.35% | 88.11% | 88.10% | 84.49% | 85.48% | 83.97% |
| Parallelization strategy efficiency | 100.00% | 99.92% | 99.50% | 98.97% | 97.44% | 99.06% | 97.78% | 97.90% | 94.26% | 95.56% | 93.94% |
| Load balancing | 100.00% | 99.94% | 99.57% | 99.10% | 97.64% | 99.37% | 98.24% | 98.38% | 94.89% | 96.48% | 94.87% |
| In execution efficiency | 100.00% | 99.98% | 99.93% | 99.87% | 99.79% | 99.69% | 99.54% | 99.51% | 99.33% | 99.04% | 99.03% |
| Scalability for computation tasks | 100.00% | 99.68% | 99.37% | 93.42% | 92.10% | 90.20% | 90.11% | 89.99% | 89.64% | 89.45% | 89.38% |
| IPC scalability | 100.00% | 99.90% | 99.78% | 99.82% | 99.78% | 99.81% | 99.85% | 99.81% | 99.72% | 99.71% | 99.67% |
| Instruction scalability | 100.00% | 100.00% | 100.00% | 100.00% | 100.00% | 99.99% | 99.99% | 99.99% | 99.99% | 99.99% | 99.99% |
| Frequency scalability | 100.00% | 99.77% | 99.60% | 93.60% | 92.30% | 90.38% | 90.25% | 90.17% | 89.90% | 89.72% | 89.69% |

Table 2: Analysis done on Wed Dec 18 04:24:29 PM CET 2024, par4115

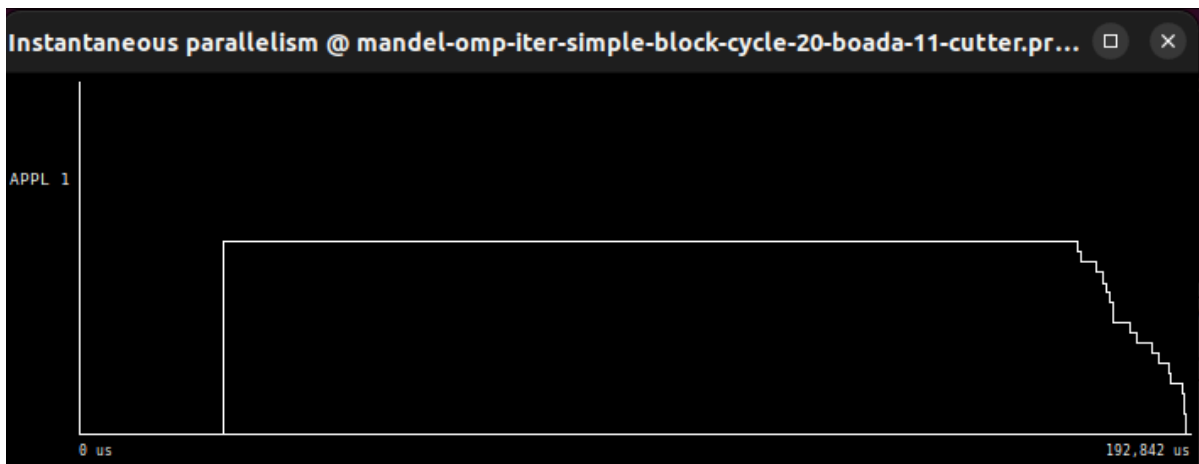
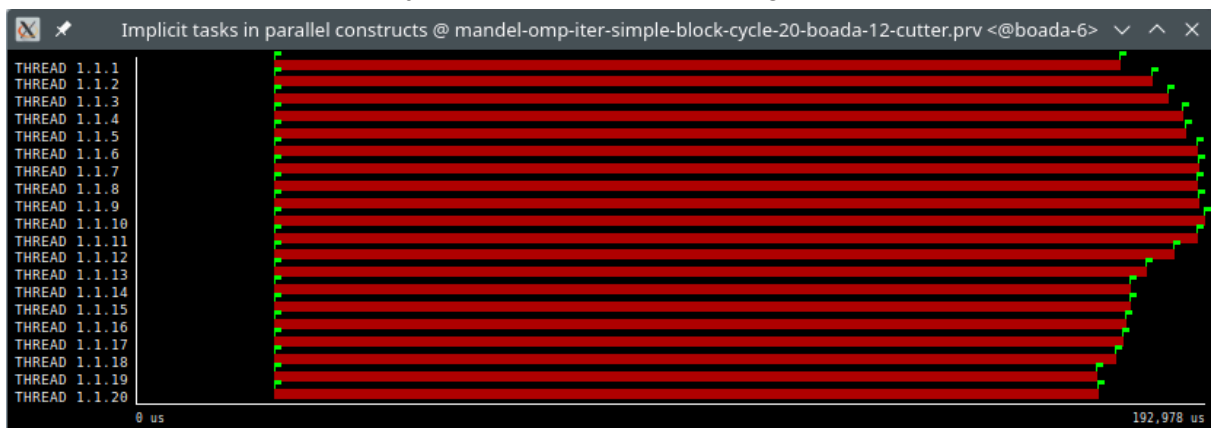
| Statistics about explicit tasks in parallel fraction | | | | | | | | | | | |
|--|------------|------------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|
| Number of processors | 1 | 2 | 4 | 6 | 8 | 10 | 12 | 14 | 16 | 18 | 20 |
| Number of implicit tasks per thread (average us) | 1.0 | 1.0 | 1.0 | 1.0 | 1.0 | 1.0 | 1.0 | 1.0 | 1.0 | 1.0 | 1.0 |
| Useful duration for implicit tasks (average us) | 2859609.21 | 1434408.88 | 719402.17 | 510166.47 | 388114.21 | 317018.05 | 264469.51 | 226983.02 | 199384.16 | 177600.92 | 159964.54 |
| Load balancing for implicit tasks | 1.0 | 1.0 | 1.0 | 0.99 | 0.98 | 0.99 | 0.98 | 0.98 | 0.95 | 0.96 | 0.95 |
| Time in synchronization implicit tasks (average us) | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 |
| Time in fork/join implicit tasks (average us) | 25.06 | 273.16 | 2956.71 | 3110.13 | 6235.34 | 1975.26 | 6960.78 | 9463.25 | 22273.89 | 6446.69 | 17526.13 |

Table 3: Analysis done on Wed Dec 18 04:24:29 PM CET 2024, par4115

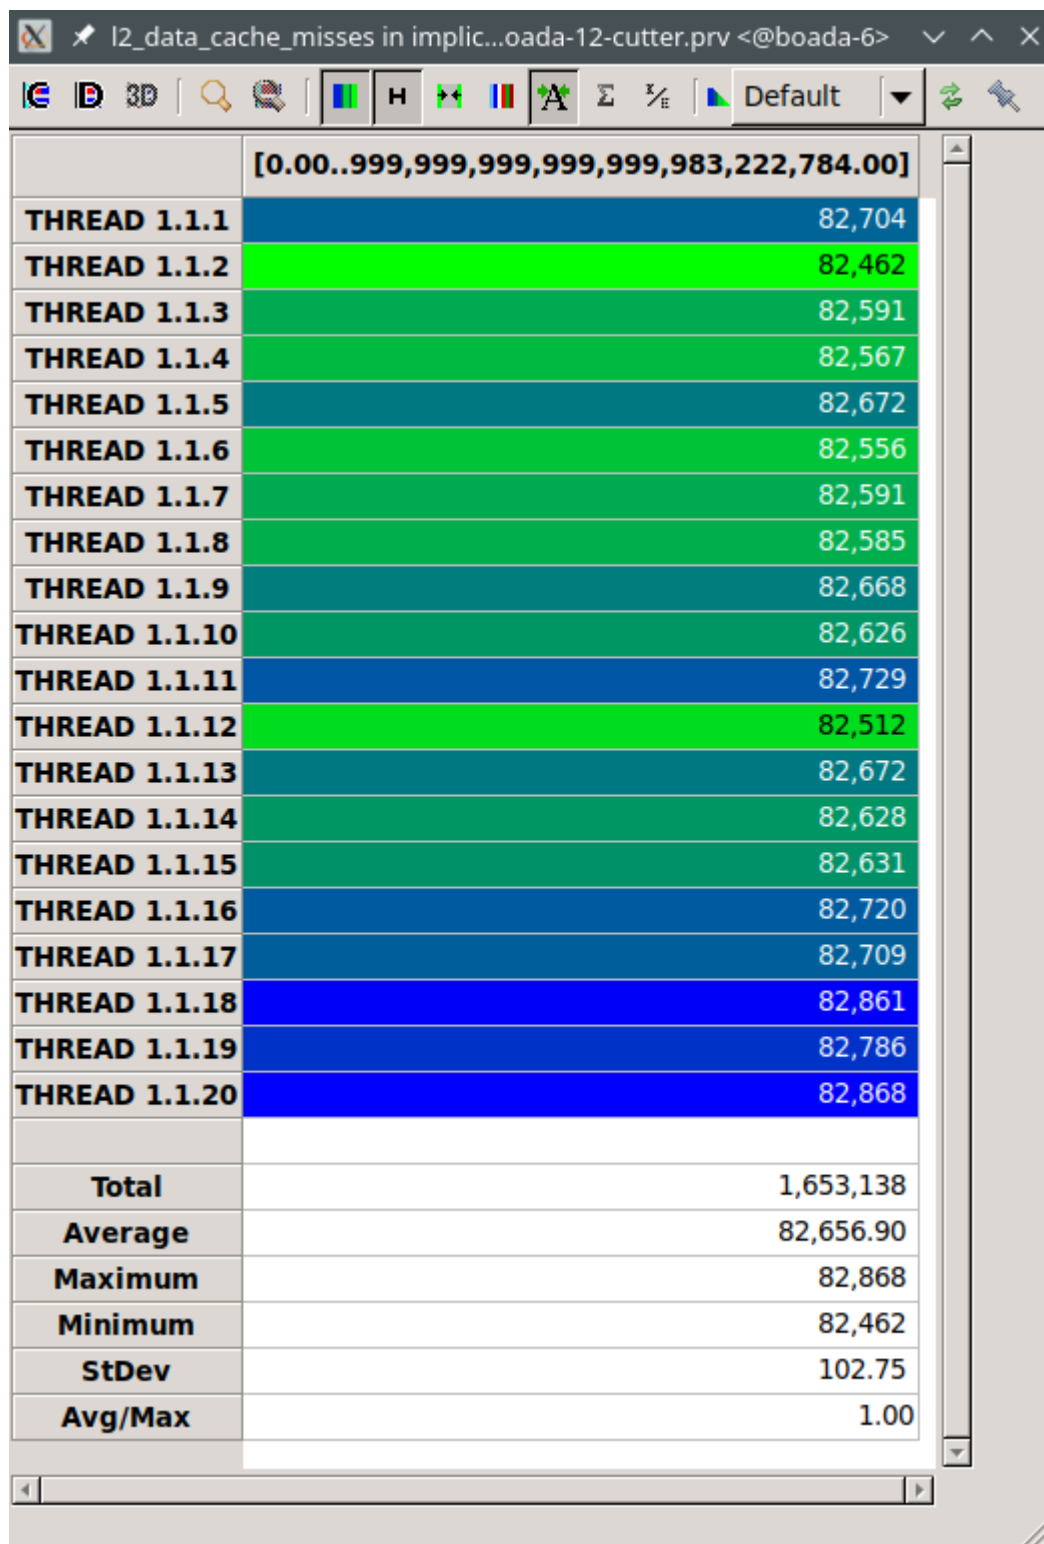
Paraver analysis



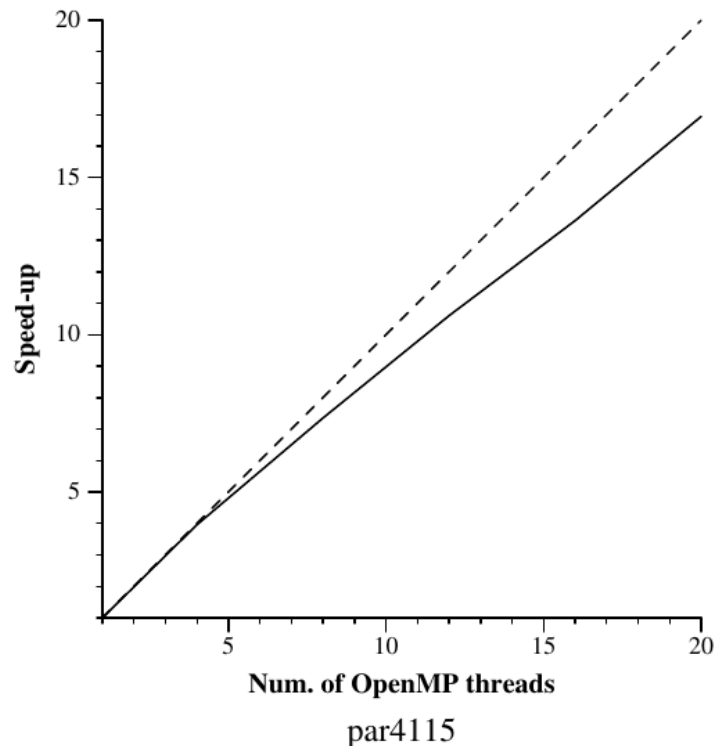
Blau ejecuta, blanc no s'ha creat, groc wait



Memory analysis



Strong scalability



Speed-up wrt sequential time (mandel function only)

Generated by par4115 on Wed Dec 18 04:16:38 PM CET 2024

Analysis

For this strategy, we can see that the model factor tables show much better numbers than the previous strategy; for 20 threads the speed-up is almost 15, the efficiency is almost 75% and the load balance almost reaches 95%. In the paraver graphics we can see how the load balance improves a lot, being almost a 100% parallelizable. The little fork & join time is explained by the few unbalances. About the misses in the L2 cache, all the threads have more or less the same amount of misses, being the maximum 82868 misses and the minimum 82462 (≈ 400 difference). The block size chosen was 16, because the size of the cache line is 64 and each element occupies 4 bytes, so $64/4 = 16$. About the strong scalability graphic, we can say that it is pretty good, although a bit could be improved.

1D Cyclic Geometric Data Decomposition by rows

Modelfactor tables

| Overview of whole program execution metrics | | | | | | | | | | | |
|---|------|------|------|------|------|------|------|-------|-------|-------|-------|
| Number of proces-sors | 1 | 2 | 4 | 6 | 8 | 10 | 12 | 14 | 16 | 18 | 20 |
| Elapsed time (sec) | 2.88 | 1.46 | 0.75 | 0.54 | 0.42 | 0.35 | 0.30 | 0.26 | 0.24 | 0.21 | 0.20 |
| Speedup | 1.00 | 1.98 | 3.86 | 5.34 | 6.92 | 8.23 | 9.70 | 10.95 | 12.01 | 13.45 | 14.48 |
| Efficiency | 1.00 | 0.99 | 0.96 | 0.89 | 0.87 | 0.82 | 0.81 | 0.78 | 0.75 | 0.75 | 0.72 |

Table 1: Analysis done on Wed Dec 18 04:36:58 PM CET 2024, par4115

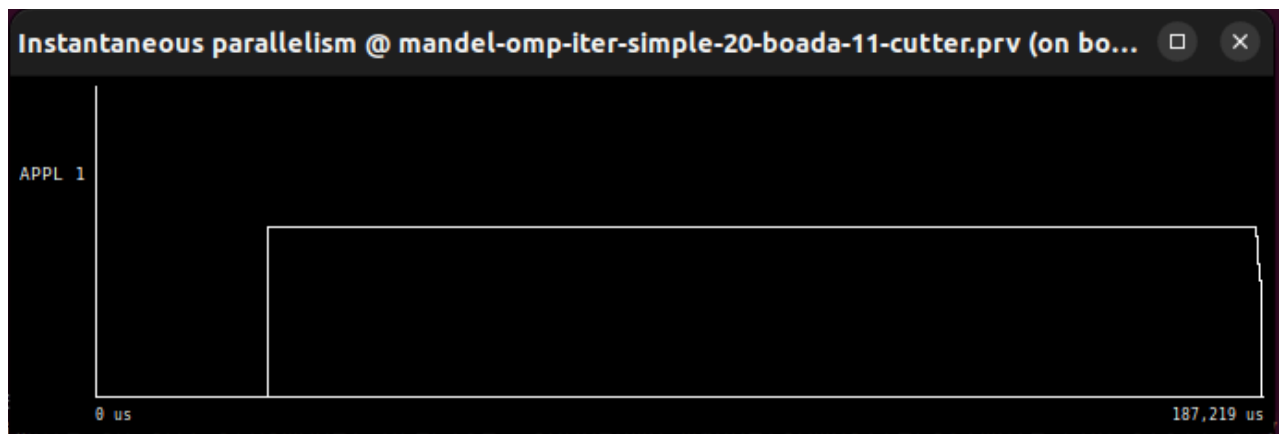
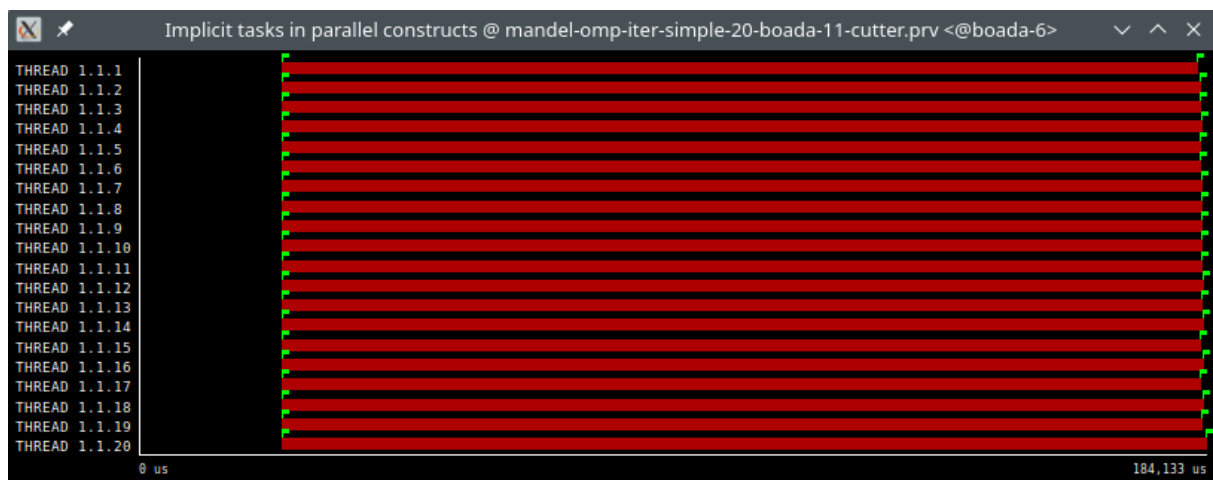

| Overview of the Efficiency metrics in parallel fraction, $\phi=99.14\%$ | | | | | | | | | | | |
|---|---------|---------|---------|---------|---------|--------|--------|--------|--------|--------|--------|
| Number of proces-sors | 1 | 2 | 4 | 6 | 8 | 10 | 12 | 14 | 16 | 18 | 20 |
| Global efficiency | 100.00% | 99.77% | 99.09% | 92.76% | 91.50% | 87.48% | 86.91% | 85.68% | 83.55% | 83.28% | 81.50% |
| Parallelization strategy efficiency | 100.00% | 99.96% | 99.82% | 99.72% | 99.29% | 99.10% | 99.05% | 98.65% | 97.27% | 98.35% | 97.90% |
| Load balancing | 100.00% | 100.00% | 99.88% | 99.85% | 99.47% | 99.41% | 99.43% | 99.25% | 97.97% | 99.11% | 98.86% |
| In execution efficiency | 100.00% | 99.96% | 99.93% | 99.87% | 99.82% | 99.69% | 99.62% | 99.40% | 99.28% | 99.24% | 99.03% |
| Scalability for computation tasks | 100.00% | 99.81% | 99.27% | 93.02% | 92.15% | 88.28% | 87.75% | 86.84% | 85.89% | 84.68% | 83.25% |
| IPC scalability | 100.00% | 99.87% | 99.68% | 99.41% | 98.64% | 97.68% | 97.19% | 96.35% | 95.49% | 94.19% | 92.94% |
| Instruction scalability | 100.00% | 100.00% | 100.00% | 100.00% | 100.00% | 99.99% | 99.99% | 99.99% | 99.99% | 99.99% | 99.99% |
| Frequency scalability | 100.00% | 99.94% | 99.59% | 93.58% | 93.42% | 90.38% | 90.29% | 90.14% | 89.96% | 89.91% | 89.59% |

Table 2: Analysis done on Wed Dec 18 04:36:58 PM CET 2024, par4115

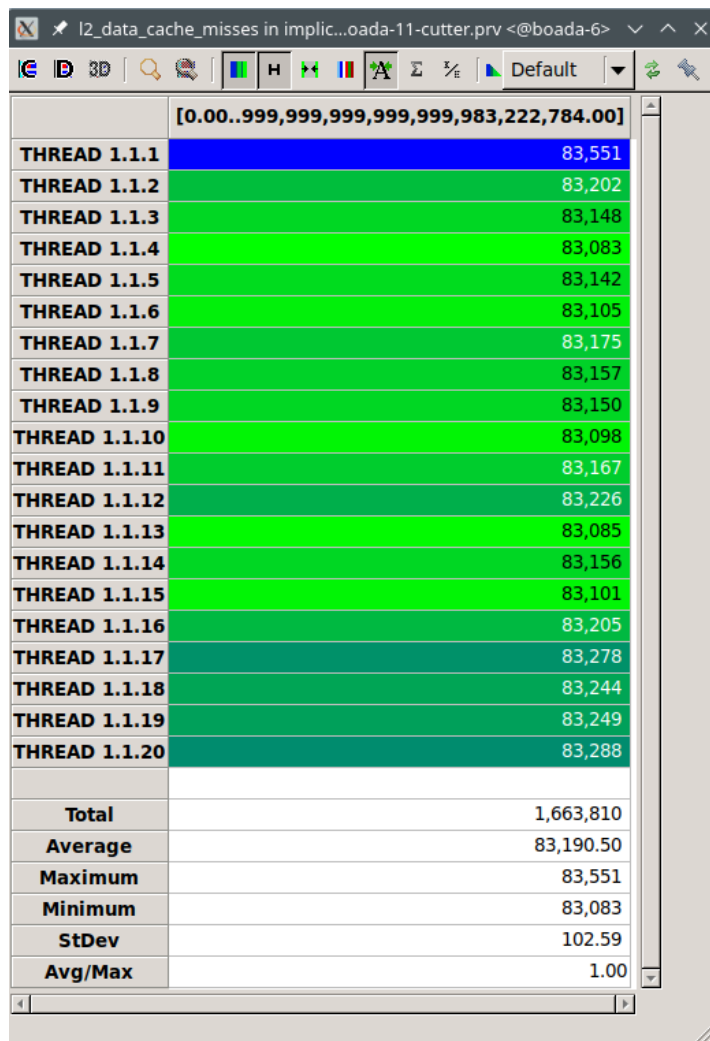
| Statistics about explicit tasks in parallel fraction | | | | | | | | | | | |
|--|------------|------------|-----------|----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|
| Number of proces-sors | 1 | 2 | 4 | 6 | 8 | 10 | 12 | 14 | 16 | 18 | 20 |
| Number of implicit tasks per thread (average us) | 1.0 | 1.0 | 1.0 | 1.0 | 1.0 | 1.0 | 1.0 | 1.0 | 1.0 | 1.0 | 1.0 |
| Useful duration for implicit tasks (average us) | 2858507.23 | 1431964.62 | 719892.19 | 512152.4 | 387755.05 | 323815.48 | 271467.67 | 235111.83 | 207996.45 | 187544.46 | 171683.58 |
| Load balancing for implicit tasks | 1.0 | 1.0 | 1.0 | 1.0 | 0.99 | 0.99 | 0.99 | 0.99 | 0.98 | 0.99 | 0.99 |
| Time in synchronization implicit tasks (average us) | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 |
| Time in fork/join implicit tasks (average us) | 117.16 | 615.67 | 1723.37 | 2136.6 | 6551.52 | 8107.53 | 8298.9 | 8931.88 | 10503.7 | 10842.22 | 12828.34 |

Table 3: Analysis done on Wed Dec 18 04:36:58 PM CET 2024, par4115


Paraver analysis



Memory analysis



Strong scalability



Speed-up wrt sequential time (mandel function only)

Generated by par4115 on Wed Dec 18 04:44:00 PM CET 2024

Analysis

The cyclic geometric data decomposition by rows is the best strategy, although it is pretty similar to the previous one. Looking at the model factor tables, we can see that the speed-up is almost 14.5, the efficiency is also good being 82% and the load balancing is almost 99% which is also shown in the paraver graphics where we can see that it has a perfect parallelization, the load balance is perfect and it is also reflected in the fork & join time, which is really low. The misses in the L2 cache differ just a tiny bit from each other and the strong scalability graphic backs it all, reaching a speed-up of almost 18 for 20 threads.