

# Parallelism (PAR)

Mastering your task decomposition strategies:  
going some steps further

Josep Ramon Herrero (T10), Gladys Utrera (T20),  
Jordi Tubella (T40), Eduard Ayguadé and Daniel Jiménez

Computer Architecture Department  
Universitat Politècnica de Catalunya

Course 2024/25 (Fall semester)

## Additional learning material for this lesson

- ▶ Atenea: Unit 3 Task decomposition
  - ▶ Video lesson 4 (overview iterative vs. recursive task decompositions) and associated questionnaire
  - ▶ Additional Atenea quizzes
  - ▶ Going further: cut-off based on number of tasks pending to be executed (optional)
- ▶ Collection of Exercises: problems in Chapter 3

# Outline

## Video lesson 4

### Task creation and synchronization (Labs summary)

### Task Decomposition

- Iterative vs Recursive

- Task generation control

- Iterative task decompositions

- Recursive task decompositions

### Reducing overheads and serialization due to synchronization

### Exploratory recursive problems

## Concepts in video lesson 4

- ▶ Linear task decomposition
  - ▶ Task = code block or procedure invocation
- ▶ (Linear) Iterative task decomposition
  - ▶ Tasks = body of iterative constructs, such as loops (countable or uncountable)
  - ▶ Examples: Pi computation, Mandelbrot in lab sessions, vector and matrix operations, ...
- ▶ Recursive task decomposition
  - ▶ Tasks = recursive procedure invocations, for example in divide-and-conquer problems
  - ▶ Examples: Fibonacci, Mandelbrot in lab sessions, graph exploration problems, ...

# Outline

Video lesson 4

Task creation and synchronization (Labs summary)

Task Decomposition

- Iterative vs Recursive

- Task generation control

- Iterative task decompositions


- Recursive task decompositions

Reducing overheads and serialization due to synchronization

Exploratory recursive problems

## Task creation in OpenMP (Labs summary)

- ▶ `#pragma omp parallel`: One **implicit** task is created for each thread in the team (and immediately executed). *There is a barrier synchronization at the end of the parallel region.*



- ▶ `int omp_get_num_threads`: returns the number of threads in the current team. 1 if outside a parallel region
- ▶ `int omp_get_thread_num`: returns the identifier of the thread in the current team that is executing a task, a value between 0 and `omp_get_num_threads()-1`

## Task creation in OpenMP (Labs summary)

- ▶ `#pragma omp single`: identifies a section of code that must be run by a single available thread. An implicit barrier exists at the end of a parallelized statement block unless the `nowait` clause is specified
- ▶ `#pragma omp task`: One **explicit** task is created, packaging code and data for (possible) deferred execution
- ▶ `#pragma omp taskloop`: **Explicit** tasks created for chunks of loop iterations. *There is a taskgroup synchronization at the end of the taskloop.*
  - ▶ In both cases, tasks executed by threads in the `parallel` region

## Task ordering in OpenMP (Labs summary)

- ▶ Thread barriers: wait for all threads to finish previous work (`#pragma omp barrier` and implicit barriers at the end of OpenMP constructs)
- ▶ Task barriers:
  - ▶ `taskwait`: Suspends the execution of the current task, waiting on the completion of its **child tasks**. The `taskwait` construct is a stand-alone directive.
  - ▶ `taskgroup`: Suspends the execution of the current task at the end of structured block, waiting on the completion of **child tasks** of the current task **and their descendent** tasks.
- ▶ Task dependences (next ...)



## taskwait vs. taskgroup

```
#pragma omp task {}          // T1
#pragma omp task             // T2
{
    #pragma omp task {}      // T3
}
#pragma omp task {}          // T4

#pragma omp taskwait
// Only T1, T2 and T4 are guaranteed to have finished at this point when T5 is created
#pragma omp task {}          // T5
```

```
#pragma omp task {}          // T1
#pragma omp taskgroup
{
    #pragma omp task          // T2
    {
        #pragma omp task {}  // T3
    }
    #pragma omp task {}       // T4
}
// Only T2, T3 and T4 are guaranteed to have finished at this point when T5 is created
#pragma omp task {}          // T5
```

# Outline

Video lesson 4

Task creation and synchronization (Labs summary)

## Task Decomposition

- Iterative vs Recursive

- Task generation control

- Iterative task decompositions

- Recursive task decompositions

Reducing overheads and serialization due to synchronization

Exploratory recursive problems



## Parallelism (PAR)

## Outline

## Video lesson 4

## Task Decomposition

## Iterative vs Recursive

## Task generation control

## Iterative task decompositions

## Recursive task decompositions

## Exploratory recursive problems

## Task generation control

Excessive task generation may not be necessary (i.e. cause excessive overhead): need mechanisms to control number of tasks and/or their granularity

- ▶ In iterative task decomposition strategies one can control task granularity by setting the number of iterations executed by each task
- ▶ In recursive task decomposition strategies one can control task granularity by controlling recursion levels where tasks are generated (**cut-off control**)
  - ▶ after certain number of recursive calls (static control)
  - ▶ when the size of the vector is too small (static control)
  - ▶ when there are sufficient tasks pending to be executed (dynamic control)
  - ▶ ...



Task granularity defined by the number of iterations out of the loop each task executes. For example, using **implicit tasks**:

Each implicit task executes a subset of iterations, based in the thread identifier executing the implicit task and the total number of implicit tasks (i.e., number of threads in the team).



Using the **work-sharing model** in OpenMP (not covered in this course):

Each implicit task executes chunks of iterations, depending on what is specified in the `schedule` clause. Implicit barrier at the end of each work-sharing (nowait clause to skip it).

## Different options to assign chunks of iterations to each implicit task through the `schedule` clause



Task granularity defined by the number of iterations each task executes. For example, using **explicit tasks**:

each explicit task executes a single iteration of the `i` loop, large task creation overhead, very fine granularity!

# Iterative task decomposition (4)

Granularity: chunk of BS loop iterations

► **Option 1:** requires loop transformation

```
void vector_add(int *A, int *B, int *C, int n) {
    int BS = ...
    for (int ii=0; ii< n; ii+=BS)
        #pragma omp task
        for (int i = ii; i < min(ii+BS, n); i++)
            C[i] = A[i] + B[i];
}

void main() {
    ...
    #pragma omp parallel
    #pragma omp single
    vector_add(a, b, c, N);
    ...
}
```

Outer loop jumps over chunks of BS iterations, inner loop traverses each chunk

## Iterative task decomposition (5)

- **Option 2:** taskloop construct to specify tasks out of loop iterations:

```
void vector_add(int *A, int *B, int *C, int n) {
    int BS = ...
    #pragma omp taskloop grainsize(BS)           // or alternatively num_tasks(n/BS)
    for (int i=0; i< n; i++)
        C[i] = A[i] + B[i];
    // Implicit task synchronization at the end of the taskloop due to the implicit taskgroup
}
void main() {
    #pragma omp parallel
    #pragma omp single
    ... vector_add(a, b, c, N); ...
}
```

- `grainsize(m)`: each task executes  $[min(m, n) .. 2 \times m]$  consecutive iterations, being  $n$  the total number of iterations
- `num_tasks(m)`: creates as many tasks as  $min(m, n)$

## Iterative task decomposition: uncountable loop

List of elements, traversed using a `while` loop while not end of list

```
int main() {
    struct node *p;

    p = init_list(n);
    ...
    #pragma omp parallel
    #pragma omp single
    while (p != NULL) {
        #pragma omp task firstprivate(p) // see note below
        process_work(p);
        p = p->next;
    }
    ...
}
```

Granularity is one iteration, hopefully with sufficient work to amortise task creation overhead.

**Note:** `firstprivate` needed to capture the value of `p` at task creation time to allow its deferred execution.

# Outline

Video lesson 4

Task creation and synchronization (Labs summary)

## Task Decomposition

- Iterative vs Recursive

- Task generation control

- Iterative task decompositions

- Recursive task decompositions

Reducing overheads and serialization due to synchronization

Exploratory recursive problems

# Recursive task decomposition: divide-and-conquer (1)

Recursively divide the problem into smaller sub-problems

```
#define N 1024
#define MIN_SIZE 64
int result = 0;

void dot_product(int *A, int *B, int n) {
    for (int i=0; i< n; i++)
        result += A[i] * B[i];
}


void rec_dot_product(int *A, int *B, int n) {
    if (n>MIN_SIZE) {
        int n2 = n / 2;
        rec_dot_product(A, B, n2);
        rec_dot_product(A+n2, B+n2, n-n2);
    }
    else
        dot_product(A, B, n);
}

void main() {
    rec_dot_product(a, b, N);
}
```




# Recursive task decomposition: divide-and-conquer (2)

N=1024, MIN\_SIZE=64



## Recursive task decomposition: leaf strategy (1)

A task corresponds with each invocation of `dot_product` once the recursive invocations stop



- Sequential generation of tasks

## Recursive task decomposition: leaf strategy (2)

```
#define N 1024
#define MIN_SIZE 64
int result = 0;

void dot_product(int *A, int *B, int n) {
    for (int i=0; i< n; i++)

        result += A[i] * B[i];
}

void rec_dot_product(int *A, int *B, int n) {
    if (n>MIN_SIZE) {
        int n2 = n / 2;
        rec_dot_product(A, B, n2);
        rec_dot_product(A+n2, B+n2, n-n2);
    }
    else
        #pragma omp task
        dot_product(A, B, n);
}

void main() {
    #pragma omp parallel
    #pragma omp single
    rec_dot_product(a, b, N);
}
```

## Recursive task decomposition: leaf strategy (3)

```
#define N 1024
#define MIN_SIZE 64
int result = 0;

void dot_product(int *A, int *B, int n) {
    for (int i=0; i< n; i++)
        #pragma omp atomic
        result += A[i] * B[i]; // one atomic per iteration
}

void rec_dot_product(int *A, int *B, int n) {
    if (n>MIN_SIZE) {
        int n2 = n / 2;
        rec_dot_product(A, B, n2);
        rec_dot_product(A+n2, B+n2, n-n2);
    }
    else
        #pragma omp task
        dot_product(A, B, n);
}
```

How could you reduce the overhead of updating variable result?

## Recursive task decomposition: leaf strategy (4)

```
#define N 1024
#define MIN_SIZE 64
int result = 0;

void dot_product(int *A, int *B, int n) {
    int tmp = 0; // local (private) variable
    for (int i=0; i< n; i++)
        tmp += A[i] * B[i];
    #pragma omp atomic
    result += tmp; // only one atomic
}

void rec_dot_product(int *A, int *B, int n) {
    if (n>MIN_SIZE) {
        int n2 = n / 2;
        rec_dot_product(A, B, n2);
        rec_dot_product(A+n2, B+n2, n-n2);
    }
    else
        #pragma omp task
        dot_product(A, B, n);
}
```

# Leaf strategy: where is the task synchronization? (1)

```
#define N 1024
#define MIN_SIZE 64
int result = 0;

void dot_product(int *A, int *B, int n);

void rec_dot_product(int *A, int *B, int n) {
    if (n>MIN_SIZE) {
        int n2 = n / 2;
        rec_dot_product(A, B, n2);
        rec_dot_product(A+n2, B+n2, n-n2);
    }
    else
        #pragma omp task
        dot_product(A, B, n);
}

void main() {
    #pragma omp parallel
    #pragma omp single
    rec_dot_product(a, b, N);
}
```

- ▶ Where is the task synchronization?
- ▶ Are there nested tasks?

## Leaf strategy: where is the task synchronization? (2)

```
#define N 1024
#define MIN_SIZE 64
int result = 0;

void dot_product(int *A, int *B, int n);

void rec_dot_product(int *A, int *B, int n) {
    if (n>MIN_SIZE) {
        int n2 = n / 2;
        rec_dot_product(A, B, n2);
        rec_dot_product(A+n2, B+n2, n-n2);
    }
    else
        #pragma omp task
        dot_product(A, B, n);
}

void main() {
    #pragma omp parallel
    #pragma omp single
    {
        rec_dot_product(a, b, N);
        // Now we need the result here.
        ....
    }
}
```

- What kind of synchronization should we use? Where?

## Leaf strategy: where is the task synchronization? (3)

```
#define N 1024
#define MIN_SIZE 64
int result = 0;

void dot_product(int *A, int *B, int n);


void rec_dot_product(int *A, int *B, int n) {
    if (n>MIN_SIZE) {
        int n2 = n / 2;
        rec_dot_product(A, B, n2);
        rec_dot_product(A+n2, B+n2, n-n2);
    }
    else
        #pragma omp task
        dot_product(A, B, n);
}

void main() {
    #pragma omp parallel
    #pragma omp single
    {
        rec_dot_product(a, b, N);
        #pragma omp taskwait
        // Now we need the result here.
        ....
    }
}
```



# How to control task granularity in leaf strategy (1)

## Leaf parallelization with **depth recursion control**




# How to control task granularity in leaf strategy (2)

## Leaf strategy with **depth recursion control**

```
#define CUTOFF 2
...
void rec_dot_product(int *A, int *B, int n, int depth) {
    if (n>MIN_SIZE) {
        int n2 = n / 2;
        if (depth == CUTOFF)
            #pragma omp task
            {
                rec_dot_product(A, B, n2, depth+1);
                rec_dot_product(A+n2, B+n2, n-n2, depth+1);
            }
        else {
            rec_dot_product(A, B, n2, depth+1);
            rec_dot_product(A+n2, B+n2, n-n2, depth+1);
        }
    }
    else // if recursion finished, need to check if task has been generated
        if (depth <= CUTOFF)
            #pragma omp task
            dot_product(A, B, n);
        else
            dot_product(A, B, n);
}
...
```

# Recursive task decomposition: tree strategy (1)

A task corresponds with each invocation of `rec_dot_product`



- ▶ Parallel generation of tasks
- ▶ Granularity: some tasks simply generate new tasks

# Recursive task decomposition: different sequential code ...

```
int dot_product(int *A, int *B, int n) {  
    int tmp = 0;  
    for (int i=0; i< n; i++) tmp += A[i] * B[i];  
    return(tmp);  
}
```

```
int rec_dot_product(int *A, int *B, int n) {  
    int tmp1, tmp2 = 0;  
    if (n>MIN_SIZE) {  
        int n2 = n / 2;  
        tmp1 = rec_dot_product(A, B, n2);  
        tmp2 = rec_dot_product(A+n2, B+n2, n-n2);  
    } else tmp1 = dot_product(A, B, n);  
    return(tmp1+tmp2);  
}
```

```
void main() {  
    result = rec_dot_product(a, b, N);  
}
```

## Recursive task decomposition: tree strategy (2)


```
int dot_product(int *A, int *B, int n) {
    int tmp = 0;
    for (int i=0; i< n; i++) tmp += A[i] * B[i];
    return(tmp);
}

int rec_dot_product(int *A, int *B, int n) {
    int tmp1, tmp2 = 0;
    if (n>MIN_SIZE) {
        int n2 = n / 2;
        #pragma omp task shared(tmp1) // firstprivate(A, B, n, n2) by default
        tmp1 = rec_dot_product(A, B, n2);
        #pragma omp task shared(tmp2) // firstprivate(A, B, n, n2) by default
        tmp2 = rec_dot_product(A+n2, B+n2, n-n2);
        #pragma omp taskwait
    } else tmp1 = dot_product(A, B, n);
    return(tmp1+tmp2);
}

void main() {
    #pragma omp parallel
    #pragma omp single
    result = rec_dot_product(a, b, N);
}
```

# How to control task granularity in tree strategy (1)

## Tree strategy with **depth recursion control**



# How to control task granularity in tree strategy (2)

## Tree strategy with **depth recursion control**

```
#define N 1024
#define MIN_SIZE 64
#define CUTOFF 3

int rec_dot_product(int *A, int *B, int n, int depth) {
    int tmp1, tmp2 = 0;
    if (n>MIN_SIZE) {
        int n2 = n / 2;
        if (depth < CUTOFF) {
            #pragma omp task shared(tmp1)
            tmp1 = rec_dot_product(A, B, n2, depth+1);
            #pragma omp task shared(tmp2)
            tmp2 = rec_dot_product(A+n2, B+n2, n-n2, depth+1);
            #pragma omp taskwait
        } else {
            tmp1 = rec_dot_product(A, B, n2, depth+1);
            tmp2 = rec_dot_product(A+n2, B+n2, n-n2, depth+1);
        }
    }
    else tmp = dot_product(A, B, n);
    return(tmp1+tmp2);
}
```

# OpenMP support for cut-off

- ▶ **final** clause: If the expression of a **final** clause evaluates to *true* the generated task and **all of its descendent tasks** will be final. The execution of a final task is sequentially **included** in the generating task (but the task is still generated)
- ▶ `omp_in_final()` intrinsic function: it returns true when executed in a final task region; otherwise, it returns false.



# OpenMP support for cut-off: tree strategy

## Making use of `omp_in_final`:

```
#define MIN_SIZE 64
#define CUTOFF 3
...
int rec_dot_product(int *A, int *B, int n, int depth) {
    int tmp1, tmp2 = 0;
    if (n>MIN_SIZE) {
        int n2 = n / 2;
        if (!omp_in_final()) {
            #pragma omp task shared(tmp1) final(depth >= CUTOFF)
            tmp1 = rec_dot_product(A, B, n2, depth+1);
            #pragma omp task shared(tmp2) final(depth >= CUTOFF)
            tmp2 = rec_dot_product(A+n2, B+n2, n-n2, depth+1);
            #pragma omp taskwait
        } else {
            tmp1 = rec_dot_product(A, B, n2, depth+1);
            tmp2 = rec_dot_product(A+n2, B+n2, n-n2, depth+1);
        }
    }
    else tmp1 = dot_product(A, B, n);
    return(tmp1+tmp2);
}
...
```



## Parallelism (PAR)

## Tree strategy: where is the task synchronization? (3)

```
int result = 0;
void dot_product(int *A, int *B, int n);

void rec_dot_product(int *A, int *B, int n) {
    if (n>MIN_SIZE) {
        int n2 = n / 2;
        #pragma omp task
        rec_dot_product(A, B, n2);
        #pragma omp task
        rec_dot_product(A+n2, B+n2, n-n2);
    } else dot_product(A, B, n);
}

void main() {
    #pragma omp parallel
    #pragma omp single
    {
        #pragma omp taskgroup
        {
            rec_dot_product(a, b, N);
        }
        // Now we need the result here..
        ....
    }
}
```

# Outline

Video lesson 4

Task creation and synchronization (Labs summary)

Task Decomposition

- Iterative vs Recursive

- Task generation control

- Iterative task decompositions

- Recursive task decompositions

Reducing overheads and serialization due to synchronization

Exploratory recursive problems

## Avoiding task barriers: task dependences (1)

- The OpenMP runtime detects dependences between sibling tasks (i.e. from the same parent task) through the specification of the directionality for the variables used in the tasks

```
#pragma omp task [depend (in : var_list)]  
                  [depend (out : var_list)]  
                  [depend (inout : var_list)]
```

Task dependences are derived from the directionality type (in, out or inout) and its items in `var_list`; this list may include array sections (e.g. `v[0:n]`)

## Avoiding task barriers: task dependences (2)

- ▶ `in` specifier: the generated task will be a dependent task of all previously generated sibling tasks that reference at least one of the list items in an `out` or `inout` list <sup>1</sup>
- ▶ `out` and `inout` specifier: the generated task will be a dependent task of all previously generated sibling tasks that reference at least one of the list items in an `in`, `out`, or `inout` list

Types of dependences:

- ▶ *read-after-write*: caused by matched `out` – `in`
- ▶ *write-after-read*: caused by matched `in` – `out`
- ▶ *write-after-write*: caused by matched `out` – `out`

---

<sup>1</sup> Note: if a list item is an array section, the matching should occur with an identically defined array section.

## Example: wavefront execution with task dependences


- ▶ Function `foo(i, j)` processes *block(i, j)*
- ▶ Wave-front execution: the execution of `foo(i, j)` depends on `foo(i-1, j)` and `foo(i, j-1)`

```
#pragma omp parallel private(i, j)
#pragma omp single
{
    for (i=1; i<n i++) {
        for (j=1; j<n;j++) {
            #pragma omp task // firstprivate(i, j) by default
                                depend(in : block[i-1][j], block[i][j-1])
                                depend(out: block[i][j])
            foo(i,j);
        }
    }
}
```



# Serialisation caused by task barriers (1)

Given a TDG to implement with the OpenMP tasking model:




```

#pragma omp task
foo1()
#pragma omp task
foo2()
#pragma omp task
foo3()
#pragma omp taskwait
#pragma omp task
foo4()
#pragma omp taskwait
#pragma omp task
foo5()
  
```


```

#pragma omp task
foo1()
#pragma omp task
foo2()
#pragma omp taskwait
#pragma omp task
foo3()
#pragma omp task
foo4()
#pragma omp taskwait
#pragma omp task
foo5()
  
```




## Serialisation caused by task barriers (2)

```
#pragma omp task
{
  #pragma omp task
  foo1()
  #pragma omp task
  foo2()
  #pragma omp taskwait
  #pragma omp task
  foo4()
  #pragma omp taskwait
}
#pragma omp task
foo3()
#pragma omp taskwait
#pragma omp task
foo5()
```



```
#pragma omp task
foo3()
#pragma omp taskgroup
{
  #pragma omp task
  foo1()
  #pragma omp task
  foo2()
}
#pragma omp task
foo4()
#pragma omp taskwait
#pragma omp task
foo5()
```



```
#pragma omp task depend(out: a)
foo1()
#pragma omp task depend(out: b)
foo2()
#pragma omp task depend(out: c)
foo3()
#pragma omp task depend(in: a, b)
depend(out: d)
foo4()
#pragma omp task depend(in: c, d)
foo5()
```



## Additional functionalities (1) (optional)

- `taskwait` with `depend` clause: instead of waiting for all child tasks to complete execution, it only waits for the predecessor child tasks according to the `in`, `out` and `inout` specifiers

```
int x=0; y=2;
```

```
#pragma omp task depend(out: x) shared(x)
compute_short1(&x);
```

```
#pragma omp task shared(y)
compute_long(&y);
```

```
#pragma omp taskwait depend(in: x)    // y not waited for at this point
printf("intermediate value for x=%d\n",x);
```

```
#pragma omp task shared(x)
compute_short2(&x)
```

```
#pragma omp taskwait
printf("final values for x=%d ; y=%d\n", x, y);
```

## Additional functionalities (2) (optional)

- ▶ An iterator can be used in the depend clause, expanding to multiple values in the specifier they appear

```
for (i = 0; i < n; ++i)
    if (i%2) {
        #pragma omp task depend(out: v[i])
        compute_element(&v[i], i);
    }


#pragma omp task depend(iterator(it = 0:n), in: v[it])
                        // could also be depend(iterator(it = 1:n:2), in: v[it])
odd = sum_odd_elements(v, n);

even = sum_even_elements(v, n);
```

Note: this is not equivalent to the use of an array section in the in specifier (i.e. `depend(in:v[0:n])`), why not?

## Additional functionalities (3) (optional)

- `mutexinoutset` specifier: equivalent to `inout` but all dependent tasks can be executed in any order, one after the other



Red dependence expressed with `depend(inout:x)` (top temporal diagram) or with `depend(mutexinoutset:x)` (bottom temporal diagram). Observe that tasks can be executed in any order, but only one at a time.

## Cancellation points in OpenMP (optional)

Tasks induced by exploratory decomposition can be terminated before finishing as soon as the desired solution is found

- ▶ `#pragma omp cancel [parallel | taskgroup]`: this directive activates the cancellation of the enclosing `[parallel | taskgroup]` region. The thread that finds the directive finishes its execution; the other threads continue their execution as normal.
- ▶ `#pragma omp cancellation point [parallel | taskgroup]`: introduces a point to check if cancellation has been activated. When found by a thread, if the enclosing `[parallel | taskgroup]` region has been already cancelled, then it finishes its execution.

## Cancellation points in OpenMP: very simple example (optional)

```
#pragma omp taskgroup
for (i=0; i<1000; i=i+100)
    #pragma omp task firstprivate(i) private(j)
    {
        for (j=i; j<i+100; j++) {
            if (do_computation(j) == 0) {
                #pragma omp cancel taskgroup
            }
            #pragma omp cancellation point taskgroup
        }
    }
```

The first task with 0 as a result of `do_computation` will finalise the execution of all the tasks in the taskgroup

# Protecting task interactions in OpenMP (Labs summary)

Two mechanisms:

1. Atomic accesses: mechanism to guarantee atomicity in load/store instructions

```
#pragma omp atomic [update | read | write]  
    expression
```

- ▶ Atomic updates: `x += 1`, `x = x - foo()`, `x[index[i]]++`
- ▶ Atomic reads: `value = *p`
- ▶ Atomic writes: `*p = value`



# Protecting task interactions in OpenMP (Labs summary)

Two mechanisms:

2. Mutual exclusion: mechanism to ensure that only one task at a time executes the code within a critical section
  - ▶ `critical` pragma: a thread waits at the beginning of a critical region until no other thread is executing a critical region (anywhere in the program)
  - ▶ `critical(name)` pragma: the `name` allows the programmer to differentiate disjoint sets of critical sections (`name` is a label, not a program variable)
  - ▶ `omp_lock_t` OpenMP intrinsics and low-level synchronization primitives (next in this chapter)

## Reducing task interactions: overhead (1)

Reductions: replicate key data structures and locally working with these local structures; when appropriate, locally replicated data structures are combined into the final global result

```
int result = 0;
// Assume this function is instantiated as a task
void dot_product(int *A, int *B, int n) {
    for (int i=0; i< n; i++)
        #pragma omp atomic
        result += A[i] * B[i];
}
```

could be easily transformed into

```
void dot_product(int *A, int *B, int n) {
    int tmp = 0;
    for (int i=0; i< n; i++)
        tmp += A[i] * B[i];

    #pragma omp atomic
    result += tmp;
}
```

## Reducing task interactions: overhead (2)

Specifying reduction operations in explicit tasks generated with either task:

```
#pragma omp parallel
#pragma omp single
{
    #pragma omp taskgroup task_reduction(+: sum)
    for (i=0; i< SIZE; i++)
        #pragma omp task firstprivate(i) in_reduction(+: sum)
        sum += X[i];
}
```

or taskloop (possible because this loop is a countable):

```
#pragma omp parallel
#pragma omp single
{
    // implicit taskgroup in taskloop construct
    #pragma omp taskloop reduction(+: sum)
    for (i=0; i< SIZE; i++)
        sum += X[i];
}
```

## Low-level synchronization functions using *locks*

**Locks:** special variables that live in memory with two basic operations:

- ▶ Acquire: while a thread has the lock, nobody else gets it; this allows the thread to do its work in private, not bothered by other threads
- ▶ Release: allow other threads to acquire the lock and do their work (one at a time) in private

Type definition and intrinsics:


```
void omp_init_lock(omp_lock_t *lock)
void omp_destroy_lock(omp_lock_t *lock)

void omp_set_lock(omp_lock_t *lock)
void omp_unset_lock(omp_lock_t *lock)

int omp_test_lock(omp_lock_t *lock)
```

## Reducing task interactions: serialization (1)

Example: inserting elements in hash table defined as a collection of linked lists



```
typedef struct {  
    int data;  
    element *next;  
} element;  
  
int dataTable[SIZE_TABLE];  
element * HashTable[SIZE_HASH];  
  
for (i = 0; i < SIZE_TABLE; i++) {  
    int index = hash_function (dataTable[i], SIZE_HASH);  
    insert_element (dataTable[i], index, HashTable);  
}
```

## Reducing task interactions: serialization (2)

Easily parallelizable using an iterative task decomposition using taskloop. However ...

- ... updates to the list in any particular slot must be protected to prevent a race condition

```
typedef struct {
    int data;
    element *next;
} element;

int dataTable[SIZE_TABLE];
element * HashTable[SIZE_HASH];

#pragma omp taskloop
for (i = 0; i < elements; i++) {
    int index = hash_function (dataTable[i], SIZE_HASH);
    #pragma omp critical // atomic not possible here
    insert_element (dataTable[i], index, HashTable);
}
```

- Serialization in the insertion of elements

## Reducing task interactions: serialization (3)

Associate a lock variable with each slot in the hash table, protecting the chain of elements in an slot

```
omp_lock_t hash_lock[SIZE_HASH];

#pragma omp parallel
#pragma omp single
{
    for (i = 0; i < SIZE_HASH; i++) omp_init_lock(&hash_lock[i]);

    #pragma omp taskloop
    for (i = 0; i < SIZE_TABLE; i++) {
        int index = hash_function (dataTable[i], SIZE_HASH);
        omp_set_lock (&hash_lock[index]);
        insert_element (dataTable[i], index, HashTable);
        omp_unset_lock (&hash_lock[index]);
    }

    for (i = 0; i < SIZE_HASH; i++) omp_destroy_lock(&hash_lock[i]);
}
```

Threads may be inserting elements into the hash table in parallel, as long as these elements hash to different slots

# Outline

Video lesson 4

Task creation and synchronization (Labs summary)

Task Decomposition

- Iterative vs Recursive

- Task generation control

- Iterative task decompositions

- Recursive task decompositions

Reducing overheads and serialization due to synchronization

Exploratory recursive problems




## How would you address the N-queens problem? (1)

```
char *a; // Solution being explored
int sol_count = 0; // Total number of solutions found
int size = 8; // board size

void nqueens(int n, int j, char *a) {
    if (j == n) sol_count += 1;
    else
        // try each possible position for queen <j>
        for ( int i=0 ; i < n ; i++ ) {
            a[j] = (char) i;
            if (ok(j + 1, a)) nqueens(n, j + 1, a);
        }
}


int main() {
    a = alloca(size * sizeof(char));
    nqueens(size, 0, a);
}
```

a = [0, 6, 3, 5, 7, 1, 4, 2]



## How would you address the N-queens problem? (2)

For a 4x4 board, the recursion tree would be ...



## How would you address the N-queens problem? (3)

```

void nqueens(int n, int j, char *a) {
    if (j == n)
        #pragma omp atomic
        sol_count += 1;
    else
        // try each possible position for queen <j>
        for ( int i=0 ; i < n ; i++ ) {
            a[j] = (char) i;
            if (ok(j + 1, a))
                #pragma omp task                                // all firstprivate by default
                nqueens(n, j + 1, a);
        }
    // Do we need to insert a task barrier at this point?
}

int main() {
    a = alloca(size * sizeof(char));
    #pragma omp parallel
    #pragma omp single
    nqueens(size, 0, a);
}

```

Do we need a new board for each task to be able to explore its own path? Is the implicit `firstprivate(a)` enough?

## How would you address the N-queens problem? (4)

A new board has to be allocated if the path is explored as a task

```
void nqueens(int n, int j, char *a) {
    if (j == n)
        #pragma omp atomic
        sol_count += 1;
    else {
        // try each possible position for queen <j>
        for ( int i=0 ; i < n ; i++ ) {
            a[j] = (char) i;
            if (ok(j + 1, a)) {
                // allocate a temporary array and copy <a> into it
                char * b = alloca(n * sizeof(char));
                memcpy(b, a, (j + 1) * sizeof(char));
                #pragma omp task                    // all firstprivate by default
                nqueens(n, j + 1, b);
            }
            #pragma omp taskwait
        }
    }
}
```

**Important:** firstprivate(b) (implicit for new board) captures the pointer to b, not the whole vector b

## Where to dynamically allocate this memory?

- ▶ `ptr=malloc(size)`: allocates memory block of given size (in bytes) in the heap, not initialized
- ▶ `ptr=alloca(size)`: as `malloc` but within the current function's stack frame; this memory will be automatically deallocated from the stack when the current function returns!

**Important:** we must insert `taskwait` if using `alloca`. For `malloc` not strictly necessary, but we have to deallocate memory


```
...
char * b = malloc(n * sizeof(char));
memcpy(b, a, (j + 1) * sizeof(char));
#pragma omp task
{
    nqueens(n, j + 1, b);
    free(b);
}
...
```

## Adding cut-off to N-queens (recursion level)

```
void nqueens(int n, int j, char *a) {
    if (j == n)
        #pragma omp atomic
        sol_count += 1;
    else
        // try each possible position for queen <j>
        if (!omp_in_final()) {
            for ( int i=0 ; i < n ; i++ ) {
                a[j] = (char) i;
                if (ok(j + 1, a))
                    // allocate a temporary array and copy <a> into it
                    char * b = alloca(n * sizeof(char));
                    memcpy(b, a, (j + 1) * sizeof(char));
                    #pragma omp task final(j>CUT_OFF)
                    nqueens(n, j + 1, b);
            }
            #pragma omp taskwait
        } else
            for ( int i=0 ; i < n ; i++ ) {
                a[j] = (char) i;
                if (ok(j + 1, a)) nqueens(n, j + 1, a);
            }
}
```

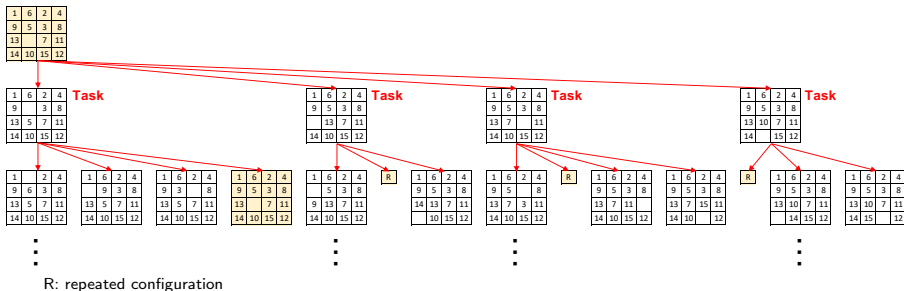
Another example: 15-puzzle (without code) ... (optional)

The solution to a 15-puzzle (a tile puzzle). Possible movements of the empty cell: UP, RIGHT, LEFT and DOWN. Here we show a series of moves that transform a given initial state to the desired final state:



## Another example: 15-puzzle (without code) ... (optional)


The state space can be explored by generating various successor states of the current state and to view them as independent tasks





## Another example: 15-puzzle (without code) ... (optional)

Anomalous speed-ups of the parallel formulation of the problem:  
the speed-up depends on where the solution is found ...



- ▶ Left:  $T_1 = 2 \times m + 1$  and  $T_4 = 1$ , therefore ...  $S_4 = 2 \times m + 1$
- ▶ Right:  $T_1 = m$  and  $T_4 = m$ , therefore ...  $S_4 = 1$

And the parallel efficiency (i.e. how well used are processors)?  
Observe that on the right three processors waste their computation

# Parallelism (PAR)

Mastering your task decomposition strategies:  
going some steps further

Josep Ramon Herrero (T10), Gladys Utrera (T20),  
Jordi Tubella (T40), Eduard Ayguadé and Daniel Jiménez

Computer Architecture Department  
Universitat Politècnica de Catalunya

Course 2024/25 (Fall semester)