


Videolesson 5 slides

# Centralized Shared Memory

- Note: System with several cores, each with its private cache memory, accessing to centralized main memory and I/O through a common bus



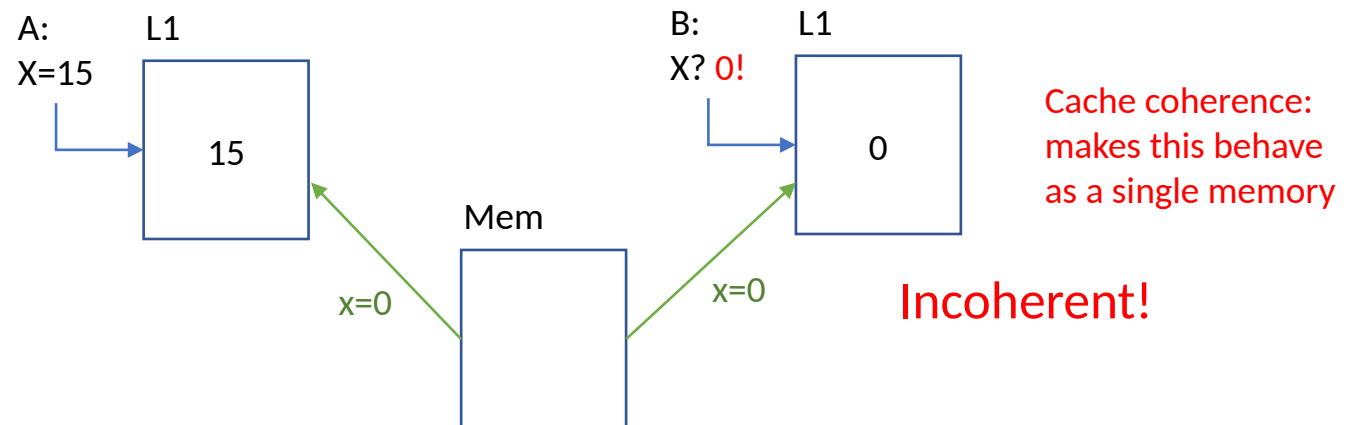
- Multi-core
- UMA: Uniform Memory Access (Time)
- SMP: Symmetric MultiProcessor

# Cache coherence problem

- Programmer: shared memory
  - Core A writes  $x=15$ , core B reads  $x$ , sees 15
- Hardware: each core has its own cache
  - One large L1 cache -> Too slow, not enough throughput
  - Private (per core) L1 caches

**Note1:** Diagram showing two cores A and B, with private cache L1, and main memory Mem.


**Note 2:** Listen to explanation in video for the specific example



# How to get coherence?

- 1) No caches (bad performance)
- 2) All cores share same L1 cache (bad performance)
- 3) Private write-through caches (Incoherent!)
- 4) Force read in one cache to see write made in another  
*on the writer side:*
  - a) Broadcast writes to update other caches (write update coherence)
  - b) Writes prevent hits to other copies (write-invalidate coherence)*on the reader side:*
  - c) Writes broadcasted on shared bus (snooping)
  - d) Each block assigned an ordering point (directory)


# Write-update snooping coherence



Note 1: diagram showing the private caches of two processors and main memory. Each cache line has a valid bit (V), tag and data. Caches access the main memory through a bus.

Note 2: Listen to explanation in video for the specific example

# Write-invalidate snooping coherence



Note 1: diagram showing the private caches of two processors and main memory. Each cache line has a valid bit (V), share bit (S) and dirty bit (D), tag and data. Caches access the main memory through a bus.

Note 2: Listen to explanation in video for the specific example