

# Contents

Implementation

Inverted files

Implementing the Boolean model

Query Optimization

Implementing the Vector model

Index compression

Getting fast the top  $r$  results



# Query answering

A **bad** algorithm:

```
input query  $q$ ;  
for every document  $d$  in database  
    check if  $d$  matches  $q$ ;  
    if so, add its docid to list  $L$ ;  
output list  $L$  (perhaps sorted in some way);
```

Time should be largely independent of database size.  
(Unavoidably) proportional to answer size.

## Central Data Structure: Inverted file

A **vocabulary** or **lexicon** or **dictionary**, usually kept in main memory, maintains all the indexed terms (*set, map...*)

- ▶ Collection: document → words contained in the document

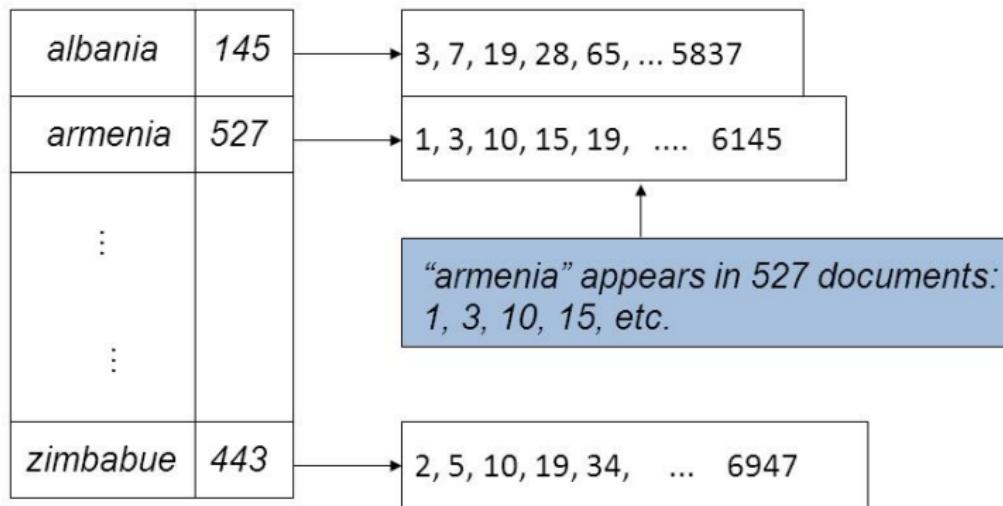
## Central Data Structure: Inverted file

A **vocabulary** or **lexicon** or **dictionary**, usually kept in main memory, maintains all the indexed terms (*set, map...*)

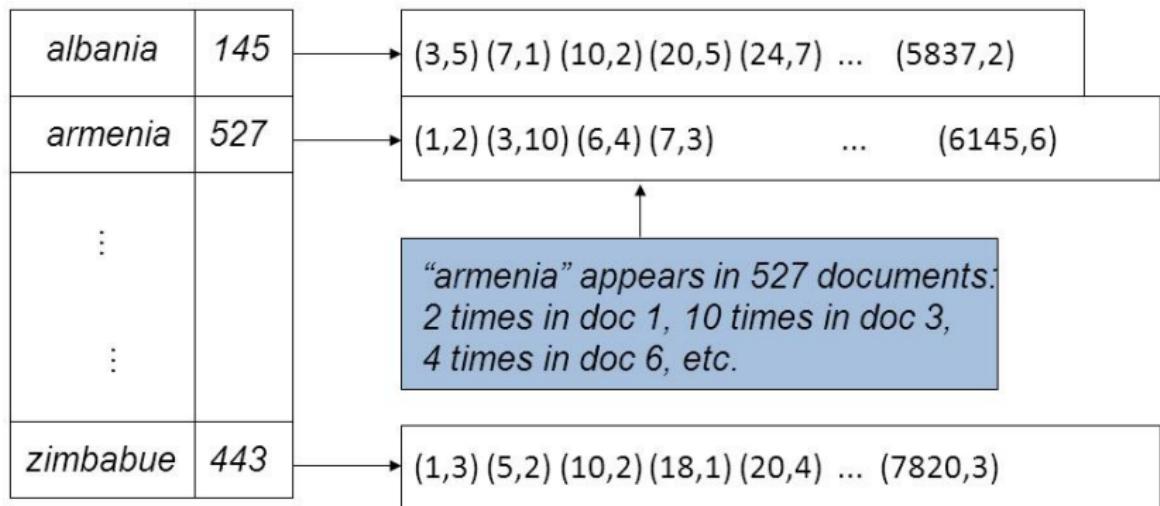
- ▶ Collection: document → words contained in the document
- ▶ Inverted file: word → documents that contain the word

Built at preprocessing time, not at query time: can afford to spend some time in its construction.

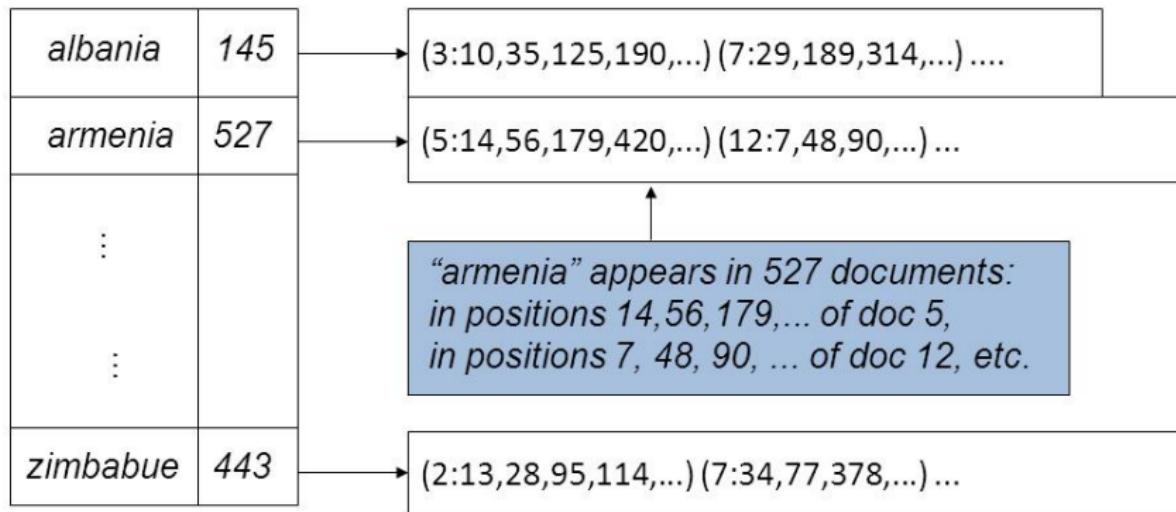
# The inverted file: Variant 1



## The inverted file: Variant 2



# The inverted file: Variant 3



# Postings

The inverted file is made of incidence/posting lists

We assign a *document identifier*, **docid** to each document.

The **dictionary** may fit in RAM for medium-size applications.

- ▶ For each indexed term, a **posting list**: list of docid's (plus maybe other info) where the term appears.
- ▶ Posting lists stored in disk for largish collections.
- ▶ Almost always sorted by *docid*.
- ▶ often compressed: minimize info to bring from disk!

# Implementation of the Boolean Model

Simplest: Traverse posting lists

Conjunctive query:  $a \text{ AND } b$

- ▶ get the **posting lists** of  $a$  and  $b$  from inverted file
- ▶ ... and intersect them
- ▶ if sorted: can do a **merge-like intersection**;
- ▶ **time**: order of the **sum** of the lengths of posting lists.

**Exercise.** Similar algorithms for OR and BUTNOT.

# Implementation of the Boolean Model

```
def intersect(L1,L2):
    i = j = 0
    Lres = []
    while i < len(L1) and j < len(L2):
        if L1[i] < L2[j]:
            ++i
        else if L1[i] > L2[j]
            ++j
        else # L1[i] == L2[j]
            Lres.append(L1[i])
            ++i
            ++j
    return Lres
```

# Query Optimization

Query Optimizer → evaluation plan for each query:

- ▶ Rewriting the query using laws of Boolean algebra
- ▶ Choosing other algorithms for intersection and union
- ▶ Using more data structures (computed offline)

# Query Rewriting

What is the most efficient way to compute  $a \text{ AND } b \text{ AND } c$ ?

- ▶  $(a \text{ AND } b) \text{ AND } c$ ?
- ▶  $(b \text{ AND } c) \text{ AND } a$ ?
- ▶  $(a \text{ AND } c) \text{ AND } b$ ?

The following are equivalent. Which is cheapest?

- ▶  $(a \text{ AND } b) \text{ OR } (a \text{ AND } c)$ ?
- ▶  $a \text{ AND } (b \text{ OR } c)$ ?

The cost of an **execution plan** depends on the sizes of the lists  
**and** the sizes of intermediate lists.

# Query Rewriting

What is the most efficient way to compute  $a \text{ AND } b \text{ AND } c$ ?

- ▶  $(a \text{ AND } b) \text{ AND } c$ ?
- ▶  $(b \text{ AND } c) \text{ AND } a$ ?
- ▶  $(a \text{ AND } c) \text{ AND } b$ ?

The following are equivalent. Which is cheapest?

- ▶  $(a \text{ AND } b) \text{ OR } (a \text{ AND } c)$ ?
- ▶  $a \text{ AND } (b \text{ OR } c)$ ?

The cost of an **execution plan** depends on the sizes of the lists  
**and** the sizes of intermediate lists.

Worst cases:

- ▶  $|L1 \cap L2| \leq \min(|L1|, |L2|)$
- ▶  $|L1 \cup L2| \leq |L1| + |L2| - |L1 \cap L2| \leq |L1| + |L2|$

# Query Rewriting

$a \text{ AND } b \text{ AND } c \rightarrow (a \text{ AND } b) \text{ AND } c$

Assume:  $|L_a| = 1.000$ ,  $|L_b| = 2.000$ ,  $|L_c| = 300$ .

Minimum comparisons if using sequential scanning =  
 $1.000 + 2.000 + 300 = 3.300$ .

# Query Rewriting

$$a \text{ AND } b \text{ AND } c \rightarrow (a \text{ AND } b) \text{ AND } c$$

Assume:  $|L_a| = 1.000$ ,  $|L_b| = 2.000$ ,  $|L_c| = 300$ .

Minimum comparisons if using sequential scanning =  
 $1.000 + 2.000 + 300 = 3.300$ .

Instruction	Comparisons	Result $\leq$
1. $L_{a \cap b} = \text{intersect}(L_a, L_b)$	$1.000 + 2.000 = 3.000$	1.000
2. $L_{res} = \text{intersect}(L_{a \cap b}, L_c)$	$1.000 + 300 = 1.300$	300
Total comparisons	$3.000 + 1.300 = 4.300$	—

# Query Rewriting

$a \text{ AND } b \text{ AND } c \rightarrow (\textcolor{red}{a \text{ AND } c}) \text{ AND } b$

Assume:  $|L_a| = 1.000$ ,  $|L_b| = 2.000$ ,  $|L_c| = 300$ .

Minimum comparisons if using sequential scanning =  
 $1.000 + 2.000 + 300 = 3.300$ .

# Query Rewriting

$$a \text{ AND } b \text{ AND } c \rightarrow (a \text{ AND } c) \text{ AND } b$$

Assume:  $|L_a| = 1.000$ ,  $|L_b| = 2.000$ ,  $|L_c| = 300$ .

Minimum comparisons if using sequential scanning =  
 $1.000 + 2.000 + 300 = 3.300$ .

Instruction	Comparisons	Result $\leq$
1. $L_{a \cap c} = \text{intersect}(L_a, L_c)$	$1.000 + 300 = 1.300$	300
2. $L_{res} = \text{intersect}(L_{a \cap c}, L_b)$	$300 + 2.000 = 2.300$	300
Total comparisons	$1.300 + 2.300 =$ <b>3.600</b> < 4.300	—

# Query Rewriting

$$a \text{ AND } b \text{ AND } c \rightarrow (a \text{ AND } c) \text{ AND } b$$

Assume:  $|L_a| = 1.000$ ,  $|L_b| = 2.000$ ,  $|L_c| = 300$ .

Minimum comparisons if using sequential scanning =  
 $1.000 + 2.000 + 300 = 3.300$ .

Instruction	Comparisons	Result $\leq$
1. $L_{a \cap c} = \text{intersect}(L_a, L_c)$	$1.000 + 300 = 1.300$	300
2. $L_{res} = \text{intersect}(L_{a \cap c}, L_b)$	$300 + 2.000 = 2.300$	300
Total comparisons	$1.300 + 2.300 =$ <b>3.600</b> < 4.300	—

Heuristic for AND-only queries: **Intersect from shortest to longest.**

# Query Rewriting

$a \text{ AND } (b \text{ OR } c)$

**Assume:**  $|L_a| = 300$ ,  $|L_b| = 4.000$ ,  $|L_c| = 5.000$ .

Minimum comparisons if using sequential scanning =  
 $300 + 4.000 + 5.000 = 9.300$ .

# Query Rewriting

$a \text{ AND } (b \text{ OR } c)$

Assume:  $|L_a| = 300$ ,  $|L_b| = 4.000$ ,  $|L_c| = 5.000$ .

Minimum comparisons if using sequential scanning =  
 $300 + 4.000 + 5.000 = 9.300$ .

Instruction	Comparisons	Result $\leq$
1. $L_{b \cup c} = \text{union}(L_b, L_c)$	$4.000 + 5.000 = 9.000$	9.000
2. $L_{res} = \text{intersect}(L_a, L_{b \cup c})$	$9.000 + 300 = 9.300$	300
Total comparisons	$9.000 + 9.300 = \mathbf{18.300}$	—

# Query Rewriting

$$a \text{ AND } (b \text{ OR } c) \rightarrow (a \text{ AND } b) \text{ OR } (a \text{ AND } c)$$

Assume:  $|L_a| = 300$ ,  $|L_b| = 4.000$ ,  $|L_c| = 5.000$ .

Minimum comparisons if using sequential scanning =  
 $300 + 4.000 + 5.000 = 9.300$ .

# Query Rewriting

$$a \text{ AND } (b \text{ OR } c) \rightarrow (a \text{ AND } b) \text{ OR } (a \text{ AND } c)$$

Assume:  $|L_a| = 300$ ,  $|L_b| = 4.000$ ,  $|L_c| = 5.000$ .

Minimum comparisons if using sequential scanning =  
 $300 + 4.000 + 5.000 = 9.300$ .

Instruction	Comparisons	Result $\leq$
1. $L_{a \cap b} = \text{intersect}(L_a, L_b)$	$300 + 4.000 = 4.300$	300
2. $L_{a \cap c} = \text{intersect}(L_a, L_c)$	$300 + 5.000 = 5.300$	300
3. $L_{res} = \text{union}(L_{a \cap b}, L_{a \cap c})$	$300 + 300 = 600$	600
Total comparisons	$4.300 + 5.300 + 600 =$ <b>9.900</b> $< 18.300$	—

# Query Rewriting

The combinatorics may get complicated . . .

$$\begin{aligned} & (a \text{ AND } b \text{ AND } d) \text{ OR } (a \text{ AND } (c \text{ OR } d) \text{ AND } e) \\ & \qquad \equiv \\ & ((a \text{ AND } d) \text{ AND } b) \text{ OR } (a \text{ AND } c \text{ AND } e) \text{ OR } ((a \text{ AND } d) \text{ AND } e) \end{aligned}$$

# Query Rewriting

The combinatorics may get complicated . . .

$$\begin{aligned} & (a \text{ AND } b \text{ AND } d) \text{ OR } (a \text{ AND } (c \text{ OR } d) \text{ AND } e) \\ & \equiv \\ & ((a \text{ AND } d) \text{ AND } b) \text{ OR } (a \text{ AND } c \text{ AND } e) \text{ OR } ((a \text{ AND } d) \text{ AND } e) \end{aligned}$$

Consider distributing so that we can compute  $\text{intersect}(L_a, L_d)$  once and store for reuse.

Exercise: Write the new plan as a sequence of instructions.  
Exercise: Find cases where the new plan is more efficient.

## Sublinear time intersection: Binary Search

Alternative: traverse one list and look up every docid in the other via **binary search**.

- ▶ **Time:** length of shortest list **times** log of length of longest.

## Sublinear time intersection: Binary Search

Alternative: traverse one list and look up every docid in the other via **binary search**.

- ▶ **Time**: length of shortest list **times** log of length of longest.

If  $|L1| \ll |L2|$

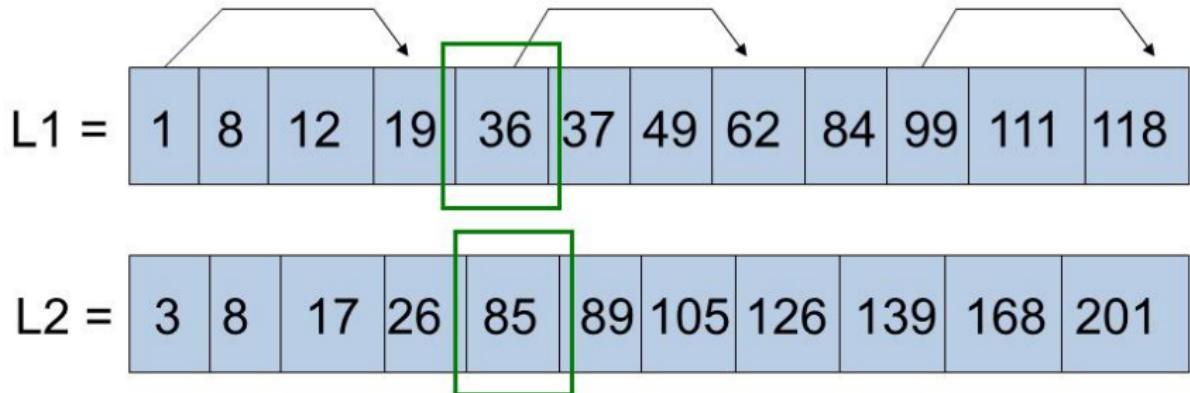
$$|L1| \cdot \log(|L2|) < |L1| + |L2|$$

# Query Optimization

Example:

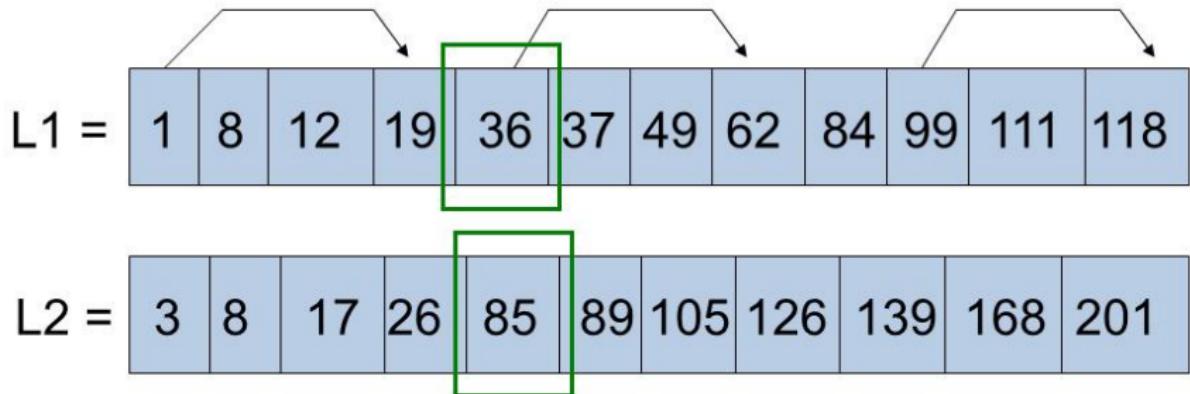
- ▶  $|L1| = 1.000, |L2| = 1.000:$ 
  - ▶ sequential scan: 2.000 comparisons,
  - ▶ binary search:  $1.000 * 10 = 10.000$  comparisons.
- ▶  $|L1| = 100, |L2| = 10.000:$ 
  - ▶ sequential scan: 10.100 comparisons,
  - ▶ binary search:  $100 * \log(10.000) = 1.400$  comparisons.

## Sublinear time intersection: Skip pointers



- ▶ We've merged 1...19 and 3...26.
- ▶ We are looking at 36 and 85.
- ▶ Since  $\text{pointer}(36)=62 < 85$ , we can jump to 84 in L1.

## Sublinear time intersection: Skip pointers



- ▶ Forward pointer from some elements.
- ▶ Either jump to next segment, or search within next segment (once).
- ▶ Optimal: in RAM,  $\sqrt{|L|}$  pointers of length  $\sqrt{|L|}$ .
- ▶ Needs random access - not so easy if in disk.

# Implementation of the Vector Model, I

## Problem statement

Fixed similarity measure  $sim(d, q)$ :

### Retrieve

documents  $d_i$  which have a similarity to the query  $q$

- ▶ either
  - ▶ above a threshold  $sim_{min}$ , or
  - ▶ the top  $r$  according to that similarity, or
  - ▶ all documents,
- ▶ sorted by decreasing similarity to the query  $q$ .

Must react **very fast** (thus, careful to the interplay with disk!), and with a reasonable memory expense.

# Implementation of the Vector Model

Obvious non-solution

```
for each d in D:  
    sim(d,q) = 0  
    get vector representing d  
    for each w in q:  
        sim(d,q) += tf(d,w) * idf(w)  
    normalize sim(d,q) by |d| * |q|  
sort results by similarity
```

$idf_w$  and  $|d|$  can be precomputed and stored in the index.  
 $|q|$  computed now.

... too inefficient for large  $D$

# Implementation of the Vector Model

Towards a faster algorithm

Most documents include a **small proportion** of the available terms.

Queries usually include a **humanly small number** of terms.

Only a very **small proportion** of the documents will be relevant.

Inverted file available!

# Implementation of the Vector Model

Idea: Invert the loops, use inverted file

```
for each w in q:  
    L = posting list for w, from inverted file  
    for each d in L:  
        if d seen for first time:  
            sim(d,q) = 0  
            sim(d,q) += tf(d,w) * idf(w)  
    for each d seen:  
        normalize sim(d,q) by |d| * |q|  
    sort results by similarity
```

# Implementation of the Vector Model

Idea: Invert the loops, use inverted file

After a few outer loops:

- ▶ Instead of having **all of**  $\text{sim}(d, q)$  for **some**  $d$ 's
- ▶ We have **partially** computed  $\text{sim}(d, q)$  for **all**  $d$ 's

= scan the document-term matrix by columns, not by rows

# Index compression, I

Why?

A large part of the query-answering time is spent  
bringing posting lists from disks to RAM.

Need to minimize amount of bits to transfer.

Index compression schemes use:

- ▶ Docid's sorted in increasing order.
- ▶ Frequencies usually very small numbers.
- ▶ Can do better than e.g. 32 bits for each.

## Index compression, II

Topic for self-study. At least:

- ▶ Unary self-delimiting code.
- ▶ Gap compression + Elias Gamma code.
- ▶ Continuation bit.
- ▶ Typical compression ratios.

E.g. books listed in the Presentation part of these notes.

## Getting fast the top $r$ results

The last line of the algorithm was

sort the documents in answer by value of  $\text{sim}(d, q)$

Time  $O(R \log R)$ , where  $R = \# \text{docs}$  with  $\text{sim}(d, q) > \text{sim}_{\min}$ .  
Noticeable if  $R$  is large (millions).

User usually wants **really fast** the top- $r$ , where  $r \ll R$ .

E.g.,  $r = 10$ .

## Example with $R = 10$ and $r = 3$

$$L = [5, 8, 3, 6, 4, 1, 10, 2, 7, 9]$$

docid	sim(docid, q)
1	0.53
2	0.26
3	0.72
4	0.32
5	0.25
6	0.36
7	0.52
8	0.88
9	0.50
10	0.19

## Getting fast the top $r$ results

Let  $L = [d_1 \dots d_R]$  be the answer (random order)

```
put [d_1, ..., d_r] in a minheap
for i = r+1..R:
    min_val = sim(d, q) for d = top of the heap
    if sim(d_i, q) > min_val:
        replace smallest element in heap with d_i
        reorganize heap
```

Claim: After any iteration, the heap contains the top- $r$  documents among the first  $i$ .

Claim: If the similarities in  $L$  are randomly ordered, the expected running time of this algorithm is  $O(R + r \cdot \ln(r) \cdot \ln(R/r))$ .

# Getting fast the top $r$ results

Let  $L = [d_1, \dots, d_R]$  be the answer

- ▶ Time to put  $r$  elements in heap:  $O(r)$ 
  - ▶ (recall why it is better than the obvious  $O(r \log r)$ )
- ▶  $\Pr(d_i \text{ enters the heap}) =$ 
$$= \Pr[d_i \text{ among } r \text{ largest in } d_1, \dots, d_i] = r/i$$
- ▶  $E[\text{time to process } d_i] = \frac{r}{i}O(\log r) + \frac{i-r}{i}O(1)$
- ▶  $E[\text{Running time}] = O(r) + \sum_{i=r+1}^R \left( \frac{r}{i}O(\log r) + \frac{i-r}{i}O(1) \right)$ 
$$= \dots \text{ (use } H(n) \simeq \ln(n), H \text{ harmonic function)}$$
$$= O(R) + O(r \ln(r) \ln(R/r))$$

For  $r \ll R$ , we go from  $O(R \log R)$  to  $O(R)$ .